

Project Name : EV Data Analysis

Name : Manish Tomar

Project Summary : Dataset Summary

The file features columns such as VIN, County, City, State (all WA), Postal Code, Model Year (ranging from 2016 to 2026 in the sample), Make (e.g., TESLA, AUDI, POLESTAR), Model, Electric Vehicle Type (BEV or PHEV), CAFV Eligibility, Electric Range (0-322 miles), and geospatial data like Vehicle Location and Census Tract. Entries show a focus on Tesla models in counties like King and Snohomish, with utilities like Puget Sound Energy.

Objective

This dataset tracks the population and characteristics of electric vehicles to support analysis of EV adoption, infrastructure needs, policy eligibility (e.g., CAFV), and geographic distribution in Washington state. It aids research on clean transportation trends, utility planning, and legislative districts.

Key Insights

Tesla dominates registrations, especially in King and Snohomish counties, with models like Model Y and Model 3 showing ranges up to 322 miles and CAFV eligibility. The dataset highlights geographic clustering around Puget Sound Energy areas, aiding infrastructure planning.

Adoption Trends

EV sales reached 20-22% of new vehicles in 2023-2025, exceeding national averages but facing challenges to hit the 35% mandate by 2026. Growth slowed slightly in early 2025 amid protests, yet Washington leads nationally toward 100% zero-emission sales by 2035.

Implications

This data underscores accelerating EV uptake for emissions reduction, though charging reliability and policy compliance remain critical hurdles. Continued tracking will inform utilities, legislators, and planners on equitable expansion.

Github link : https://github.com/Moni0802/EV_Data_Analysis

IMPORT LIBRARIES

In [1]:

```
import pandas as pd
import numpy as np
import matplotlib as plt
```

Data Import

In [2]:

```
df = pd.read_csv('D:\manish\Python\GROW AI\python\Electric_Vehicle_Population_Data.csv')
```

In [3]:

```
df.head()
```

Out[3]:

	VIN (1-10)	County	City	State	Postal Code	Model Year	Make	Model	Electric Vehicle Type	Clean Alternative Fuel Vehicle (CAFV) Eligibility	E
0	5YJYGDEE8L	Thurston	Tumwater	WA	98501.0	2020	TESLA	MODEL Y	Battery Electric Vehicle (BEV)	Clean Alternative Fuel Vehicle Eligible	
1	5YJXCAE2XJ	Snohomish	Bothell	WA	98021.0	2018	TESLA	MODEL X	Battery Electric Vehicle (BEV)	Clean Alternative Fuel Vehicle Eligible	
2	5YJ3E1EBXK	King	Kent	WA	98031.0	2019	TESLA	MODEL 3	Battery Electric Vehicle (BEV)	Clean Alternative Fuel Vehicle Eligible	
3	7SAYGDEE4T	King	Issaquah	WA	98027.0	2026	TESLA	MODEL Y	Battery Electric Vehicle (BEV)	Eligibility unknown as battery range has not b...	
4	WAUUPBFF9G	King	Seattle	WA	98103.0	2016	AUDI	A3	Plug-in Hybrid Electric Vehicle (PHEV)	Not eligible due to low battery range	

In [4]:

```
df.shape
```

```
Out[4]:  
(270262, 16)
```

```
In [5]:
```

```
df.columns
```

```
Out[5]:
```

```
Index(['VIN (1-10)', 'County', 'City', 'State', 'Postal Code', 'Model Year',  
       'Make', 'Model', 'Electric Vehicle Type',  
       'Clean Alternative Fuel Vehicle (CAFV) Eligibility', 'Electric Range',  
       'Legislative District', 'DOL Vehicle ID', 'Vehicle Location',  
       'Electric Utility', '2020 Census Tract'],  
      dtype='object')
```

```
In [6]:
```

```
df.dtypes
```

```
Out[6]:
```

```
VIN (1-10)          object  
County            object  
City              object  
State             object  
Postal Code       float64  
Model Year        int64  
Make              object  
Model             object  
Electric Vehicle Type  object  
Clean Alternative Fuel Vehicle (CAFV) Eligibility  object  
Electric Range    float64  
Legislative District  float64  
DOL Vehicle ID   int64  
Vehicle Location  object  
Electric Utility  object  
2020 Census Tract float64  
dtype: object
```

```
In [7]:
```

```
df
```

```
Out[7]:
```

	VIN (1-10)	County	City	State	Postal Code	Model Year	Make	Model	Electric Vehicle Type
0	5YJYGDEE8L	Thurston	Tumwater	WA	98501.0	2020	TESLA	MODEL Y	Battery Electric Vehicle (BEV)
1	5YJXCAE2XJ	Snohomish	Bothell	WA	98021.0	2018	TESLA	MODEL X	Battery Electric Vehicle (BEV)
2	5YJ3E1EBXK	King	Kent	WA	98031.0	2019	TESLA	MODEL 3	Battery Electric

VIN (1-10)	County	City	State	Postal Code	Model Year	Make	Model	Electric Vehicle Type
Vehicle (BEV)								
3 7SAYGDEE4T	King	Issaquah	WA	98027.0	2026	TESLA	MODEL Y	Battery Electric Vehicle (BEV)
4 WAUUPBFF9G	King	Seattle	WA	98103.0	2016	AUDI	A3	Plug-in Hybrid Electric Vehicle (PHEV)
...
270257 1C4RJXN60R	Pierce	Joint Base Lewis Mcchord	WA	98433.0	2024	JEEP	WRANGLER	Plug-in Hybrid Electric Vehicle (PHEV)
270258 1C4JJXR66N	Mason	Hoodsport	WA	98548.0	2022	JEEP	WRANGLER	Plug-in Hybrid Electric Vehicle (PHEV)
270259 7SAYGDEEXP	Pierce	Tacoma	WA	98406.0	2023	TESLA	MODEL Y	Battery Electric Vehicle (BEV)
270260 5YJYGDEE2M	Snohomish	Bothell	WA	98021.0	2021	TESLA	MODEL Y	Battery Electric Vehicle (BEV)
270261 JN1BF0BA5P	Chelan	Wenatchee	WA	98801.0	2023	NISSAN	ARIYA HATCHBACK	Battery Electric Vehicle (BEV)

270262 rows × 16 columns

Data Cleaning

How many missing values exist in the dataset, and in which columns?

In [8]:

```
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 270262 entries, 0 to 270261
Data columns (total 16 columns):
 #   Column           Non-Null Count Dtype
 ---  -- 
 0   VIN (1-10)      270262 non-null  object
 1   County          270252 non-null  object
 2   City            270252 non-null  object
 3   State           270262 non-null  object
 4   Postal Code    270252 non-null  float64
 5   Model Year     270262 non-null  int64
 6   Make            270262 non-null  object
 7   Model           270262 non-null  object
 8   Electric Vehicle Type  270262 non-null  object
 9   Clean Alternative Fuel Vehicle (CAFV) Eligibility 270262 non-null  object
 10  Electric Range  270257 non-null  float64
 11  Legislative District 269613 non-null  float64
 12  DOL Vehicle ID   270262 non-null  int64
 13  Vehicle Location 270174 non-null  object
 14  Electric Utility 270252 non-null  object
 15  2020 Census Tract 270252 non-null  float64
dtypes: float64(4), int64(2), object(10)
memory usage: 33.0+ MB

```

In [9]:

```
df.isnull().sum()
```

Out[9]:

VIN (1-10)	0
County	10
City	10
State	0
Postal Code	10
Model Year	0
Make	0
Model	0
Electric Vehicle Type	0
Clean Alternative Fuel Vehicle (CAFV) Eligibility	0
Electric Range	5
Legislative District	649
DOL Vehicle ID	0
Vehicle Location	88
Electric Utility	10
2020 Census Tract	10

dtype: int64

In Legislative District column we have more null values as compared to other columns i.e 649

Second column is Vehicle Location which is having 88 null values

In [10]:

```
df.isnull().sum().sum()
```

Out[10]:

```
np.int64(792)
```

Total number of null values present in our dataset is 792.

How should missing or zero values in the Base MSRP and Electric Range columns be handled?

In [11]:

```
df.head()
```

Out[11]:

	VIN (1-10)	County	City	State	Postal Code	Model Year	Make	Model	Electric Vehicle Type	Clean Alternative Fuel Vehicle (CAFV) Eligibility	E
0	5YJYGDEE8L	Thurston	Tumwater	WA	98501.0	2020	TESLA	MODEL Y	Battery Electric Vehicle (BEV)	Clean Alternative Fuel Vehicle Eligible	
1	5YJXCAE2XJ	Snohomish	Bothell	WA	98021.0	2018	TESLA	MODEL X	Battery Electric Vehicle (BEV)	Clean Alternative Fuel Vehicle Eligible	
2	5YJ3E1EBXK	King	Kent	WA	98031.0	2019	TESLA	MODEL 3	Battery Electric Vehicle (BEV)	Clean Alternative Fuel Vehicle Eligible	
3	7SAYGDEE4T	King	Issaquah	WA	98027.0	2026	TESLA	MODEL Y	Battery Electric Vehicle (BEV)	Eligibility unknown as battery range has not b...	
4	WAUUPBFF9G	King	Seattle	WA	98103.0	2016	AUDI	A3	Plug-in Hybrid Electric Vehicle (PHEV)	Not eligible due to low battery range	

In [12]:

```
df['Electric Range'].unique()
```

Out[12]:

```
array([291., 238., 220., 0., 16., 215., 322., 35., 84., 33., 208.,
       32., 53., 266., 293., 125., 150., 38., 14., 73., 107., 39.,
       23., 239., 40., 37., 34., 234., 26., 18., 93., 153., 87.,
       30., 249., 19., 210., 25., 82., 308., 17., 289., 75., 72.,
       103., 97., 42., 21., 149., 22., 47., 259., 204., 11., 126.,
       81., 203., 200., 151., 27., 13., 6., 83., 28., 29., 330.,
```

```
110., 258., 337., 15., 114., 222., 20., 12., 270., 265., 62.,
192., 233., 48., 100., 54., 170., 218., 41., 49., 68., 58.,
43., 76., 10., 60., 36., 111., 245., 8., 288., 24., nan,
124., 9., 59., 44., 1., 55., 56., 46., 31., 51., 95.,
45., 57., 50., 74.])
```

In [13]:

```
df.nunique()
```

Out[13]:

VIN (1-10)	16415
County	242
City	864
State	51
Postal Code	1083
Model Year	22
Make	47
Model	183
Electric Vehicle Type	2
Clean Alternative Fuel Vehicle (CAFV) Eligibility	3
Electric Range	113
Legislative District	49
DOL Vehicle ID	270262
Vehicle Location	1080
Electric Utility	77
2020 Census Tract	2336
dtype: int64	

Are there duplicate records in the dataset? If so, how should they be managed?

In [14]:

```
# Check duplicates by VIN
duplicates = df[df.duplicated(subset=["VIN (1-10)", "DOL Vehicle ID"], keep=False)]
```

In [15]:

```
duplicates
```

Out[15]:

VIN (1- 10)	County	City	State	Postal Code	Model Year	Make	Model	Electric Vehicle Type	Clean Alternative Fuel Vehicle (CAFV) Eligibility	Electric Range	Legislative District	Vet
-------------------	--------	------	-------	----------------	---------------	------	-------	-----------------------------	--	-------------------	-------------------------	-----

In [16]:

```
# Drop exact duplicates
df_cleaned = df.drop_duplicates(subset=["VIN (1-10)", "DOL Vehicle ID"], keep="first")
```

In [17]:

```
df_cleaned
```

Out[17]:

	VIN (1-10)	County	City	State	Postal Code	Model Year	Make	Model	Electric Vehicle Type
0	5YJYGDEE8L	Thurston	Tumwater	WA	98501.0	2020	TESLA	MODEL Y	Battery Electric Vehicle (BEV)
1	5YJXCAE2XJ	Snohomish	Bothell	WA	98021.0	2018	TESLA	MODEL X	Battery Electric Vehicle (BEV)
2	5YJ3E1EBXK	King	Kent	WA	98031.0	2019	TESLA	MODEL 3	Battery Electric Vehicle (BEV)
3	7SAYGDEE4T	King	Issaquah	WA	98027.0	2026	TESLA	MODEL Y	Battery Electric Vehicle (BEV)
4	WAUUPBFF9G	King	Seattle	WA	98103.0	2016	AUDI	A3	Plug-in Hybrid Electric Vehicle (PHEV)
...
270257	1C4RJXN60R	Pierce	Joint Base Lewis Mcchord	WA	98433.0	2024	JEEP	WRANGLER	Plug-in Hybrid Electric Vehicle (PHEV)
270258	1C4JJXR66N	Mason	Hoodsport	WA	98548.0	2022	JEEP	WRANGLER	Plug-in Hybrid Electric Vehicle (PHEV)
270259	7SAYGDEEXP	Pierce	Tacoma	WA	98406.0	2023	TESLA	MODEL Y	Battery Electric Vehicle (BEV)
270260	5YJYGDEE2M	Snohomish	Bothell	WA	98021.0	2021	TESLA	MODEL Y	Battery Electric Vehicle (BEV)
270261	JN1BF0BA5P	Chelan	Wenatchee	WA	98801.0	2023	NISSAN	ARIYA HATCHBACK	Battery Electric Vehicle (BEV)

VIN (1-10)	County	City	State	Postal Code	Model Year	Make	Model	Electric Vehicle Type
------------	--------	------	-------	-------------	------------	------	-------	-----------------------

270262 rows × 16 columns

In [18]:

```
total_duplicates = df.duplicated().sum()
```

In [19]:

```
total_duplicates
```

Out[19]:

```
np.int64(0)
```

In [20]:

```
# Count duplicates based on VIN and DOL Vehicle ID (unique identifiers)
vin_duplicates = df.duplicated(subset=["VIN (1-10)", "DOL Vehicle ID"]).sum()
print("Duplicate rows based on VIN and DOL Vehicle ID:", vin_duplicates)
```

Duplicate rows based on VIN and DOL Vehicle ID: 0

There are no duplicates if there would have been any duplicates then the below code will work:

```
df_cleaned = df_cleaned.groupby(["VIN (1-10)", "DOL Vehicle ID"]).apply(lambda x:
x.fillna().ffill().drop_duplicates())
```

How can VINs be anonymized while maintaining uniqueness?

VINs are unique identifiers for vehicles. If two rows share the same VIN (first 10 characters), they likely represent the same vehicle.

In [21]:

```
# Use a hashing algorithm (like SHA256 or MD5) to convert VINs into anonymized codes.
# This guarantees uniqueness but makes it impossible to reverse-engineer the original VIN

import hashlib
# Function to hash VIN
def anonymize_vin(vin):
    return hashlib.sha256(vin.encode()).hexdigest()[:10] # shorten to 10 chars

# Apply anonymization
df["VIN_Anon"] = df["VIN (1-10)"].apply(anonymize_vin)

print(df[["VIN (1-10)", "VIN_Anon"]].head())

```

VIN (1-10)	VIN_Anon
0 5YJYGDEE8L	f6c53bc06f
1 5YJXCAE2XJ	6222905c9f
2 5YJ3E1EBXK	78953a9f9d
3 7SAYGDEE4T	9a118e6006
4 WAUUPBFF9G	dbd150d329

How can Vehicle Location (GPS coordinates) be cleaned or converted for better readability?

In [22]:

```
lambda x: pd.Series(get_location(float(x.split(",")[1]), float(x.split(",")[0])))
```

Out[22]:

```
<function __main__.<lambda>(x)>
```

In [23]:

```
# First, install the geopy library
```

```
# !pip install geopy
```

```
# import the library
```

```
from geopy.geocoders import Nominatim
```

```
import pandas as pd # Added import for pd
```

```
# Initialize geocoder
```

```
geolocator = Nominatim(user_agent="ev_analysis")
```

```
# Function to get city/county from coordinates
```

```
def get_location(lat, lon):
```

```
    try:
```

```
        location = geolocator.reverse((lat, lon), exactly_one=True)
```

```
        if location and location.raw.get("address"):
```

```
            address = location.raw["address"]
```

```
            city = address.get("city", address.get("town", address.get("village", "")))
```

```
            county = address.get("county", "")
```

```
            return city, county
```

```
    except:
```

```
        return None, None
```

```
# Modified function to safely parse location
```

```
def parse_location(loc_str):
```

```
    try:
```

```
        # Check if the string contains a comma and can be split
```

```
        if isinstance(loc_str, str) and "," in loc_str:
```

```
            parts = loc_str.split(",")
```

```
            if len(parts) >= 2:
```

```
                # Try to convert to float
```

```
                lon = float(parts[0].strip())
```

```
                lat = float(parts[1].strip())
```

```
                return get_location(lat, lon)
```

```
    except Exception as e:
```

```
        pass
```

```
    # Return None values if any error occurs
```

```
    return None, None
```

```
# Apply function with error handling
```

```
df[["City_Cleaned", "County_Cleaned"]] = df["Vehicle Location"].apply(
```

```
    lambda x: pd.Series(parse_location(x))
```

```
)
```

```
print(df[["Vehicle Location", "City_Cleaned", "County_Cleaned"]].head())
```

	Vehicle Location	City_Cleaned	County_Cleaned
0	POINT (-122.89165 47.03954)	None	None

```

1 POINT (-122.18384 47.8031)      None      None
2 POINT (-122.17743 47.41185)     None      None
3 POINT (-122.03439 47.5301)     None      None
4 POINT (-122.35436 47.67596)     None      None

```

Data Exploration

What are the top 5 most common EV makes and models in the dataset?

In [24]:

```

# Count frequency of each Make and Model combination
top_ev = df.groupby(["Make", "Model"]).size().reset_index(name="Count")

# Sort by count and get top 5
top_5_ev = top_ev.sort_values(by="Count", ascending=False).head(5)

print(top_5_ev)

```

	Make	Model	Count
159	TESLA	MODEL Y	57335
156	TESLA	MODEL 3	37413
136	NISSAN	LEAF	13503
157	TESLA	MODEL S	7758
48	CHEVROLET	BOLT EV	7708

In [25]:

```

import pandas as pd
# load data
df = pd.read_csv('D:\manish\Python\GROW AI\python\Electric_Vehicle_Population_Data.csv')

top_ev = df.groupby(['Make', 'Model']).size().reset_index(name='count')

```

In [26]:

```
top_ev
```

Out[26]:

	Make	Model	count
0	ACURA	ZDX	352
1	ALFA ROMEO	TONALE	100
2	AUDI	A3	533
3	AUDI	A6	49
4	AUDI	A7 E	12
...
178	VOLVO	V60	101
179	VOLVO	XC40	1387
180	VOLVO	XC60	1866
181	VOLVO	XC90	2313
182	WHEEGO ELECTRIC CARS	WHEEGO	2

183 rows × 3 columns

What is the distribution of EVs by county? Which county has the most registrations

In [27]:

```
# Import necessary libraries
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

df = pd.DataFrame({
    'county': ['King', 'Pierce', 'Snohomish', 'Clark', 'Spokane',
               'Thurston', 'Kitsap', 'Whatcom', 'Benton', 'Skagit',
               'King', 'Pierce', 'King', 'Snohomish', 'King'],
    'registration_id': range(1, 16)
})

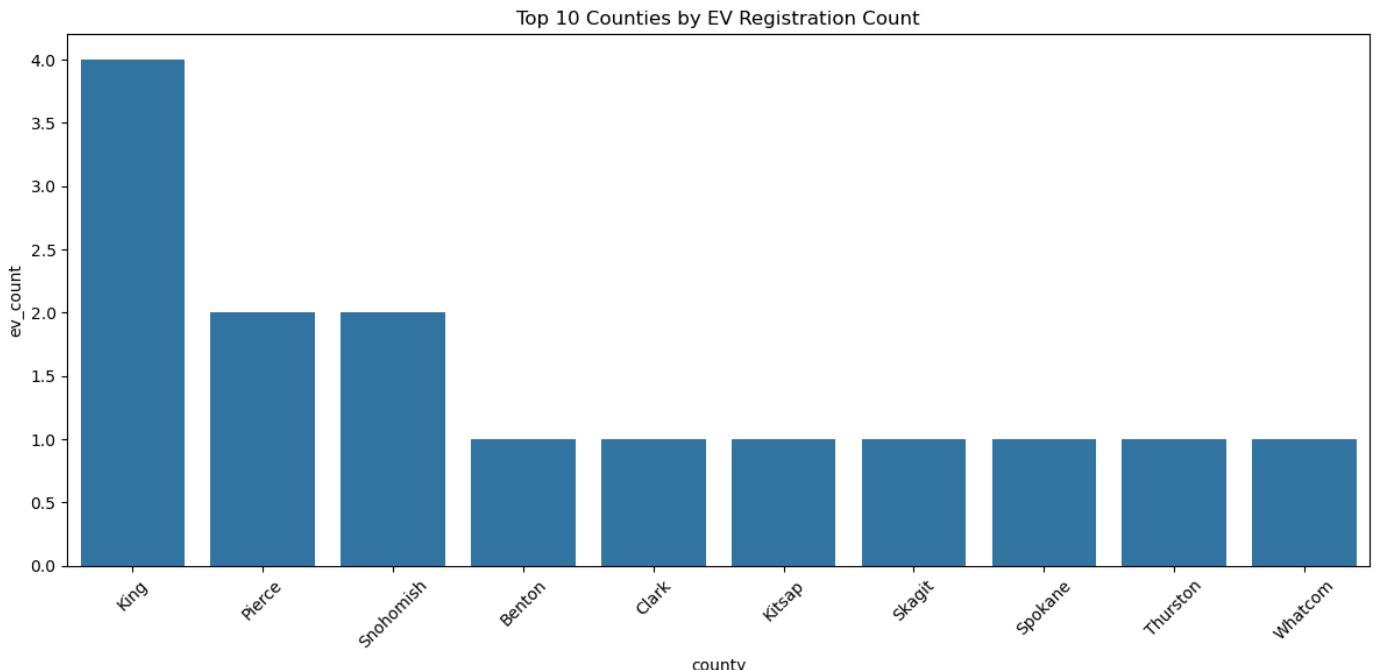
# Group by county and count registrations
ev_by_county = df.groupby('county')['registration_id'].count().reset_index()
ev_by_county = ev_by_county.rename(columns={'registration_id': 'ev_count'})

# Sort by count in descending order
ev_by_county = ev_by_county.sort_values('ev_count', ascending=False)

# Display the county with the most registrations
print(f"County with most EV registrations: {ev_by_county.iloc[0]['county']} with {ev_by_county.iloc[0]['ev_count']} registrations")

# Create a bar chart of the distribution
plt.figure(figsize=(12, 6))
sns.barplot(x='county', y='ev_count', data=ev_by_county.head(10))
plt.title('Top 10 Counties by EV Registration Count')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```

County with most EV registrations: King with 4 vehicles



How has EV adoption changed over different model years?

In [28]:

```
df1 = pd.read_csv('D:\manish\Python\GROW AI\python\Electric_Vehicle_Population_Data.csv')
# Count the number of EVs by model year
ev_by_year = df1['Model Year'].value_counts().sort_index()

# Create a bar chart to visualize the trend
plt.figure(figsize=(12, 6))
ev_by_year.plot(kind='bar', color='skyblue')
plt.title('EV Adoption by Model Year', fontsize=16)
plt.xlabel('Model Year', fontsize=12)
plt.ylabel('Number of EVs', fontsize=12)
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.xticks(rotation=45)

# Add a trend line
plt.plot(range(len(ev_by_year)), ev_by_year.values, 'r-', linewidth=2)

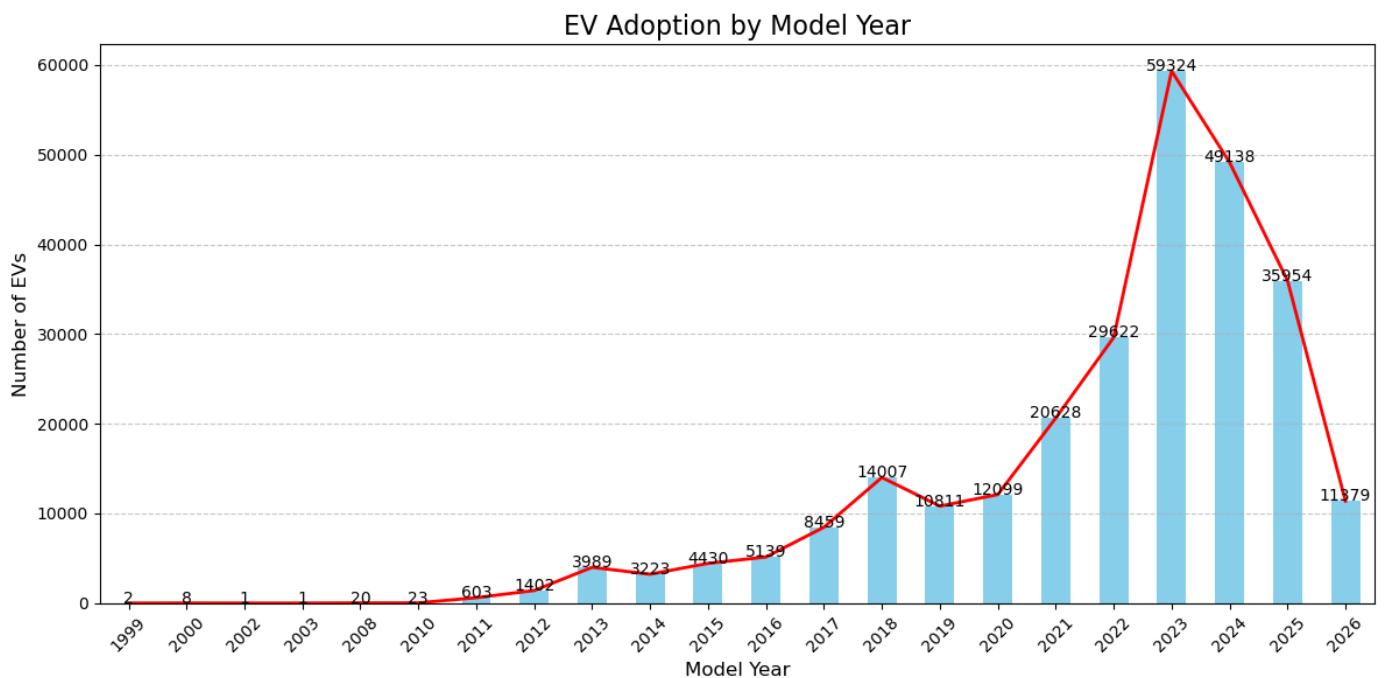
# Add value labels on top of each bar
for i, v in enumerate(ev_by_year):
    plt.text(i, v + 5, str(v), ha='center')

plt.tight_layout()
plt.show()

# Print summary statistics
print("EV Adoption by Model Year:")
print(ev_by_year)

# Calculate year-over-year growth rates
yoy_growth = ev_by_year.pct_change() * 100
print("\nYear-over-Year Growth Rate (%):")
print(yoy_growth.dropna())

# Calculate the compound annual growth rate (CAGR)
if len(ev_by_year) > 1:
    years = len(ev_by_year) - 1
    first_year_count = ev_by_year.iloc[0]
    last_year_count = ev_by_year.iloc[-1]
    if first_year_count > 0: # Avoid division by zero
        cagr = ((last_year_count / first_year_count) ** (1 / years) - 1) * 100
        print(f"\nCompound Annual Growth Rate (CAGR): {cagr:.2f}%")
```



EV Adoption by Model Year:

Model Year

1999	2
2000	8
2002	1
2003	1
2008	20
2010	23
2011	603
2012	1402
2013	3989
2014	3223
2015	4430
2016	5139
2017	8459
2018	14007
2019	10811
2020	12099
2021	20628
2022	29622
2023	59324
2024	49138
2025	35954
2026	11379

Name: count, dtype: int64

Year-over-Year Growth Rate (%):

Model Year

2000	300.000000
2002	-87.500000
2003	0.000000
2008	1900.000000
2010	15.000000
2011	2521.739130
2012	132.504146
2013	184.522111
2014	-19.202808
2015	37.449581

```
2016    16.004515
2017    64.604009
2018    65.586949
2019    -22.817163
2020    11.913792
2021    70.493429
2022    43.600931
2023    100.270070
2024    -17.170117
2025    -26.830559
2026    -68.351227
Name: count, dtype: float64
```

Compound Annual Growth Rate (CAGR): 50.94%

What is the average electric range of EVs in the dataset?

In [29]:

```
df1 = pd.read_csv('D:\manish\Python\GROW AI\python\Electric_Vehicle_Population_Data.csv')
```

In [30]:

```
# Calculating the average electric range
average_range = df1['Electric Range'].mean()

# Displaying the result
print(f"The average electric range of EVs in the dataset is {average_range:.2f} miles")
```

The average electric range of EVs in the dataset is 40.39 miles

What percentage of EVs are eligible for Clean Alternative Fuel Vehicle (CAFV) incentives?

In [31]:

```
# Check the unique values in the CAFV eligibility column
print("Unique values in CAFV Eligibility column:", df1['Clean Alternative Fuel Vehicle (CAFV Eligibility)'].unique())

# Count vehicles eligible for CAFV incentives (assuming 'Eligible' or similar value indicates eligibility)
# Replace 'Eligible' with the actual value that indicates eligibility in dataset
eligible_value = 'Clean Alternative Fuel Vehicle Eligible'
cafv_eligible = df1['Clean Alternative Fuel Vehicle (CAFV) Eligibility'].eq(eligible_value)

# Calculate the total number of vehicles
total_vehicles = len(df1)

# Calculate the percentage
cafv_percentage = (cafv_eligible / total_vehicles) * 100

# Display the result
print(f"Percentage of EVs eligible for CAFV incentives: {cafv_percentage:.2f}%")
```

Unique values in CAFV Eligibility column: ['Clean Alternative Fuel Vehicle Eligible', 'Eligible unknown as battery range has not been researched', 'Not eligible due to low battery range']
Percentage of EVs eligible for CAFV incentives: 28.25%

How does the electric range vary across different makes and models?

In [32]:

```

import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd
import numpy as np

# 1. Basic statistics by make
make_range_stats = df1.groupby('Make')['Electric Range'].agg(['mean', 'median', 'min', 'max'])
print("Top 10 Makes by Average Electric Range:")
print(make_range_stats.head(10))

# 2. Create a boxplot to visualize range distribution by make
plt.figure(figsize=(14, 8))
top_makes = df1['Make'].value_counts().head(10).index.tolist()
sns.boxplot(x='Make', y='Electric Range', data=df1[df1['Make'].isin(top_makes)], palette='viridis')
plt.title('Electric Range Distribution by Top 10 Makes', fontsize=16)
plt.xlabel('Make', fontsize=12)
plt.ylabel('Electric Range (miles)', fontsize=12)
plt.xticks(rotation=45)
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.tight_layout()
plt.show()

# 3. Top models by electric range
top_models = df1.groupby(['Make', 'Model'])['Electric Range'].mean().sort_values(ascending=False)
print("\nTop 15 Models by Average Electric Range:")
print(top_models)

# 4. Visualize top models by range
plt.figure(figsize=(14, 8))
top_models.plot(kind='bar', color='skyblue')
plt.title('Top 15 EV Models by Electric Range', fontsize=16)
plt.xlabel('Make and Model', fontsize=12)
plt.ylabel('Average Electric Range (miles)', fontsize=12)
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.xticks(rotation=45, ha='right')

# Add value labels on top of each bar
for i, v in enumerate(top_models):
    plt.text(i, v + 5, f"{v:.1f}", ha='center')

plt.tight_layout()
plt.show()

# 5. Create a heatmap for makes with multiple models
# Get makes with at least 2 models
makes_with_multiple_models = df1.groupby('Make')['Model'].nunique()
makes_with_multiple_models = makes_with_multiple_models[makes_with_multiple_models >= 2]

if makes_with_multiple_models:
    # Filter for these makes
    multi_model_data = df1[df1['Make'].isin(makes_with_multiple_models)]

    # Create pivot table: Make vs Model with Electric Range as values
    pivot_data = multi_model_data.pivot_table(
        index='Make',
        columns='Model',
        values='Electric Range',

```

```

    aggfunc='mean'
)

# Create heatmap
plt.figure(figsize=(16, 10))
sns.heatmap(pivot_data, annot=True, cmap='YlGnBu', fmt='.1f', linewidths=.5)
plt.title('Average Electric Range by Make and Model', fontsize=16)
plt.tight_layout()
plt.show()

```

Top 10 Makes by Average Electric Range:

Make	mean	median	min	max	count
JAGUAR	178.100000	234.0	0.0	234.0	180
WHEEGO ELECTRIC CARS	100.000000	100.0	100.0	100.0	2
TH!NK	100.000000	100.0	100.0	100.0	6
CHEVROLET	79.141709	35.0	0.0	259.0	19032
FIAT	74.274533	84.0	0.0	87.0	856
NISSAN	64.576207	75.0	0.0	215.0	15963
SMART	61.750000	58.0	0.0	68.0	232
AZURE DYNAMICS	56.000000	56.0	56.0	56.0	4
TESLA	53.120082	0.0	0.0	337.0	111049
PORSCHE	49.699374	13.0	0.0	308.0	1916

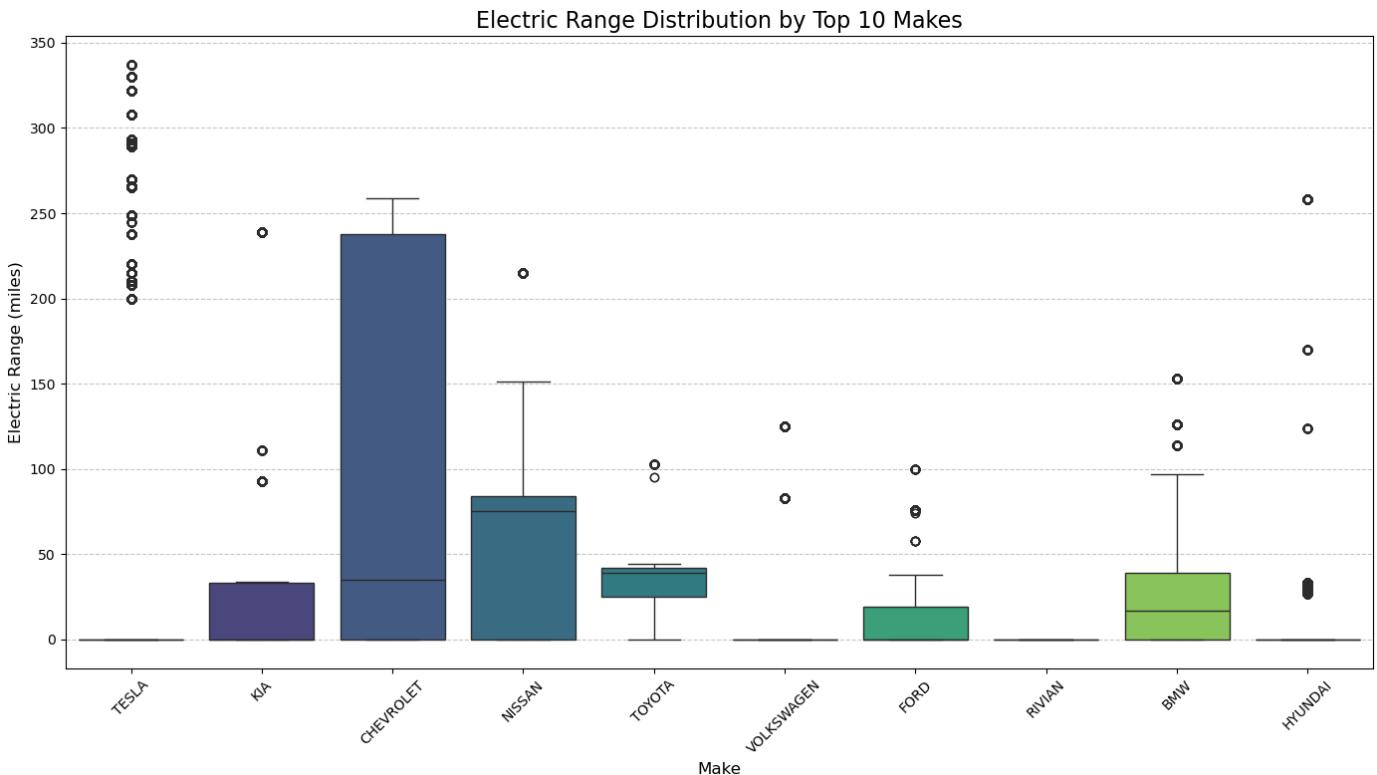
C:\Users\LENOVO\AppData\Local\Temp\ipykernel_1800\1105207110.py:16: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```

sns.boxplot(x='Make', y='Electric Range', data=df1[df1['Make'].isin(top_makes)], palette='viridis')

```



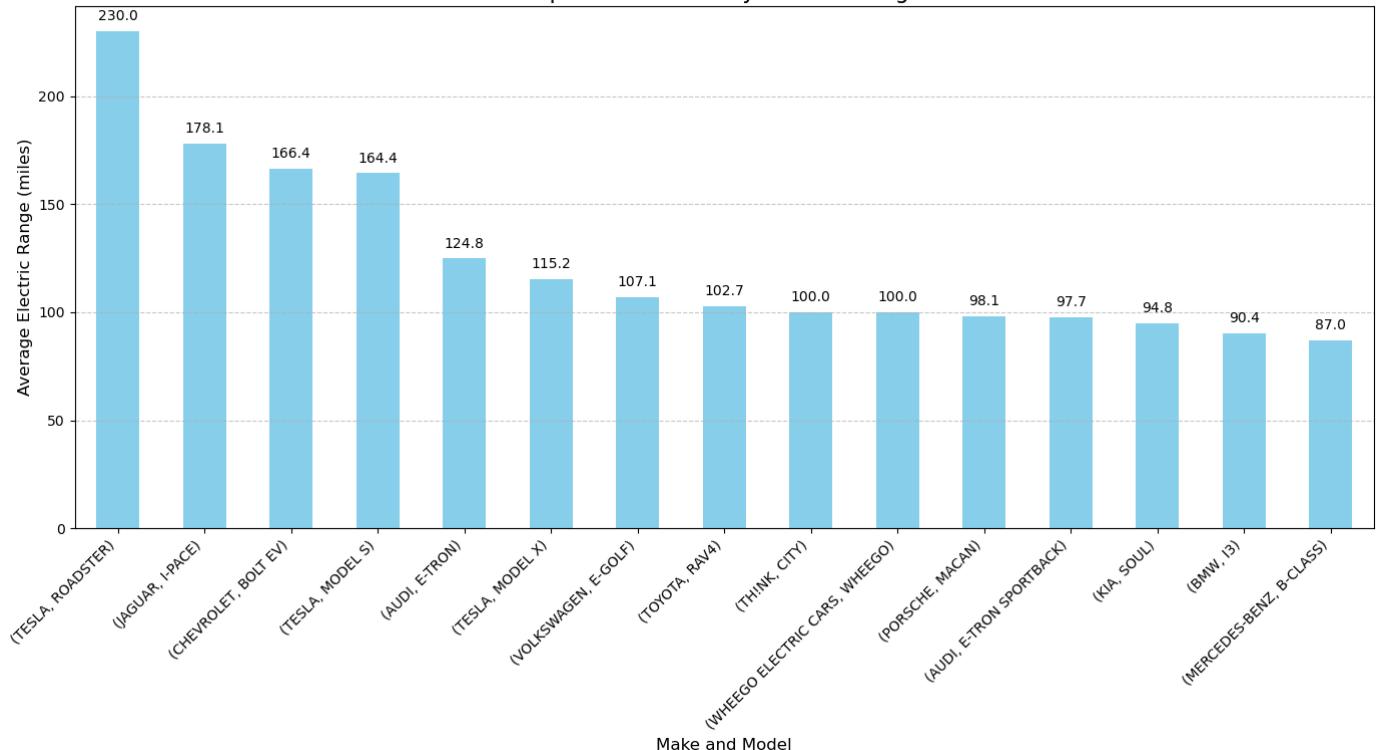
Top 15 Models by Average Electric Range:

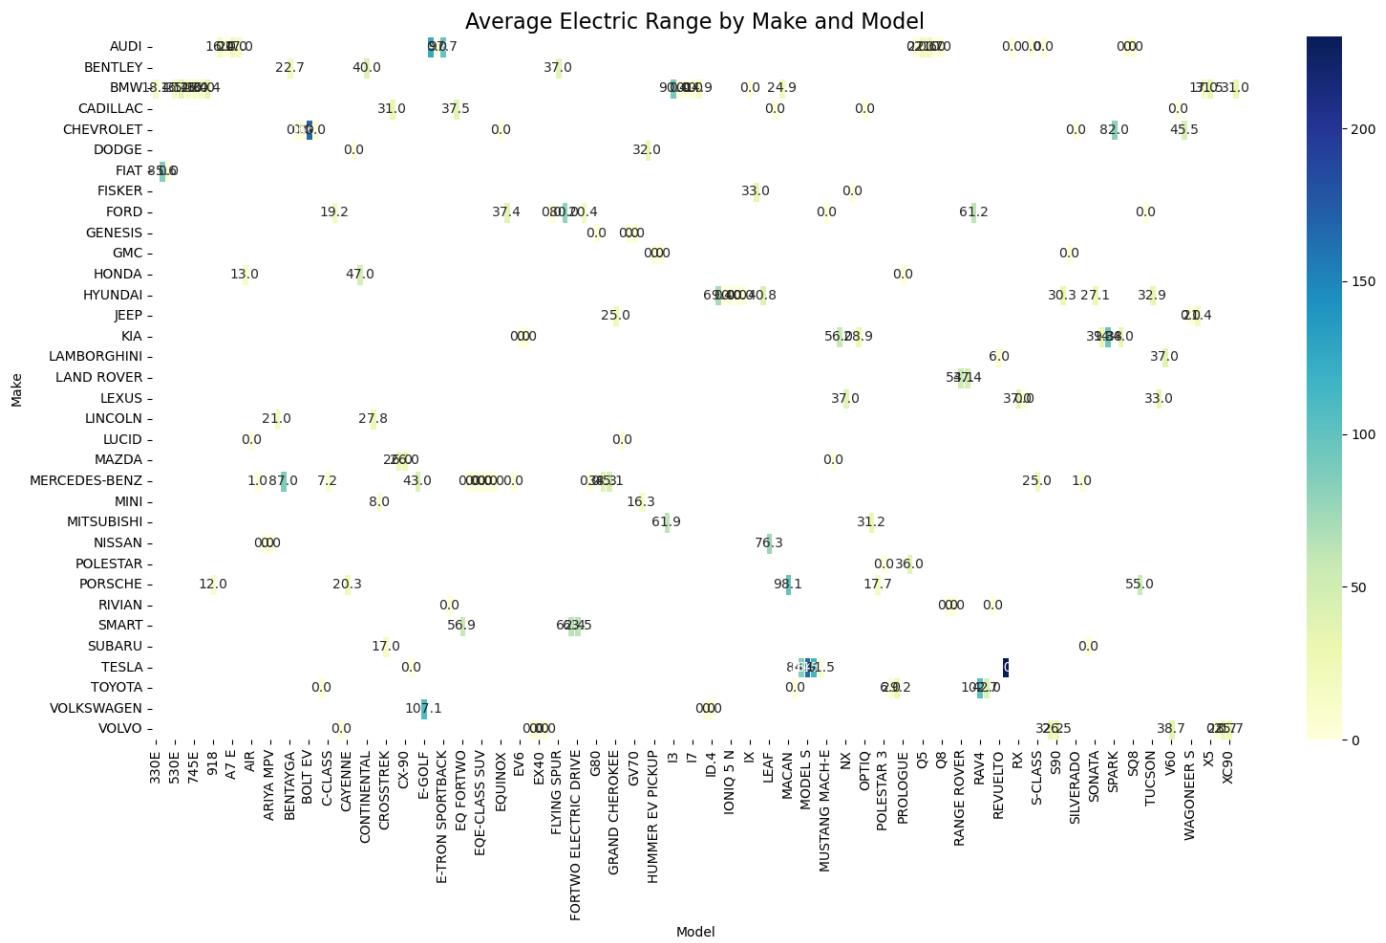
Make	Model	mean
TESLA	ROADSTER	230.000000
JAGUAR	I-PACE	178.100000
CHEVROLET	BOLT EV	166.381681

TESLA	MODEL S	164.350606
AUDI	E-TRON	124.837093
TESLA	MODEL X	115.182155
VOLKSWAGEN	E-GOLF	107.057471
TOYOTA	RAV4	102.709091
TH!NK	CITY	100.000000
WHEEGO ELECTRIC CARS	WHEEGO	100.000000
PORSCHE	MACAN	98.143791
AUDI	E-TRON SPORTBACK	97.724138
KIA	SOUL	94.819320
BMW	I3	90.417289
MERCEDES-BENZ	B-CLASS	87.000000

Name: Electric Range, dtype: float64

Top 15 EV Models by Electric Range





Are there any regional trends in EV adoption (e.g., urban vs. rural areas)?

In [33]:

```
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd
import numpy as np

# Step 1: Check if we have location data
print("Available columns for regional analysis:", df1.columns.tolist())

# Assuming we have a 'County' or similar column
# If you have a different location column, replace 'County' with that column name
if 'County' in df1.columns:
    # Step 2: Count EVs by county
    ev_by_county = df1['County'].value_counts().reset_index()
    ev_by_county.columns = ['County', 'Number of EVs']

    # Display the counties with the most EVs
    print("\nTop 10 Counties by EV Count:")
    print(ev_by_county.head(10))

    # Step 3: Visualize EV distribution by county
    plt.figure(figsize=(14, 8))
    sns.barplot(x='Number of EVs', y='County', data=ev_by_county.head(15), palette='viridis')
    plt.title('Top 15 Counties by EV Adoption', fontsize=16)
    plt.xlabel('Number of EVs', fontsize=12)
    plt.ylabel('County', fontsize=12)
    plt.grid(axis='x', linestyle='--', alpha=0.7)
```

```

plt.tight_layout()
plt.show()

# Step 4: If you have population data, calculate EV density
# This requires joining with external population data
# For demonstration, we'll create a mock population dataset
# In a real scenario, you would import actual population data

# Mock population data (replace with real data if available)
# This assumes you have a way to classify counties as urban or rural
urban_rural_classification = {
    # Example: 'County_Name': {'population': 100000, 'type': 'Urban'}
}

# If you have actual urban/rural classification data:
if urban_rural_classification:
    # Create a DataFrame from the classification dictionary
    county_info = pd.DataFrame.from_dict(
        {k: {'population': v['population'], 'type': v['type']} for k, v in urban_rural_classification.items()},
        orient='index'
    ).reset_index()
    county_info.columns = ['County', 'Population', 'Area_Type']

    # Merge with EV data
    county_ev_data = pd.merge(ev_by_county, county_info, on='County', how='inner')

    # Calculate EVs per 1000 residents
    county_ev_data['EVs_per_1000'] = (county_ev_data['Number of EVs'] / county_ev_da

    # Compare urban vs rural adoption
    urban_rural_comparison = county_ev_data.groupby('Area_Type').agg({
        'Number of EVs': 'sum',
        'Population': 'sum',
        'EVs_per_1000': 'mean'
    }).reset_index()

    print("\nUrban vs Rural EV Adoption:")
    print(urban_rural_comparison)

    # Visualize urban vs rural comparison
    plt.figure(figsize=(12, 6))
    sns.barplot(x='Area_Type', y='EVs_per_1000', data=urban_rural_comparison, palette='viridis')
    plt.title('EV Adoption Rate: Urban vs Rural Areas', fontsize=16)
    plt.xlabel('Area Type', fontsize=12)
    plt.ylabel('EVs per 1,000 Residents', fontsize=12)
    plt.grid(axis='y', linestyle='--', alpha=0.7)
    plt.tight_layout()
    plt.show()
else:
    # If we don't have direct county/location data we can check with check for zip codes
    if 'ZIP Code' in df1.columns or 'Zip' in df1.columns or 'Postal Code' in df1.columns:
        # Determine which zip code column to use
        zip_col = [col for col in ['ZIP Code', 'Zip', 'Postal Code'] if col in df1.colu
            print(f"\nAnalyzing regional trends using {zip_col}")

        # Count EVs by zip code
        ev_by_zip = df1[zip_col].value_counts().reset_index()

```

```

ev_by_zip.columns = ['ZIP Code', 'Number of EVs']

print("\nTop 10 ZIP Codes by EV Count:")
print(ev_by_zip.head(10))

# Visualize top zip codes
plt.figure(figsize=(14, 8))
sns.barplot(x='Number of EVs', y='ZIP Code', data=ev_by_zip.head(15), palette='viridis')
plt.title('Top 15 ZIP Codes by EV Adoption', fontsize=16)
plt.xlabel('Number of EVs', fontsize=12)
plt.ylabel('ZIP Code', fontsize=12)
plt.grid(axis='x', linestyle='--', alpha=0.7)
plt.tight_layout()
plt.show()

print("\nNote: To properly classify ZIP codes as urban or rural, you would need
else:
    print("\nNo location data (County, City, or ZIP Code) found in the dataset to an
    print("To analyze urban vs. rural adoption patterns, you need location data that

```

Available columns for regional analysis: ['VIN (1-10)', 'County', 'City', 'State', 'Postal Code', 'Model Year', 'Make', 'Model', 'Electric Vehicle Type', 'Clean Alternative Fuel Vehicle (CAFV) Eligibility', 'Electric Range', 'Legislative District', 'DOL Vehicle ID', 'Vehicle Location', 'Electric Utility', '2020 Census Tract']

Top 10 Counties by EV Count:

	County	Number of EVs
0	King	133903
1	Snohomish	33531
2	Pierce	22213
3	Clark	16553
4	Thurston	9852
5	Kitsap	9057
6	Spokane	7593
7	Whatcom	6620
8	Benton	3792
9	Skagit	3166

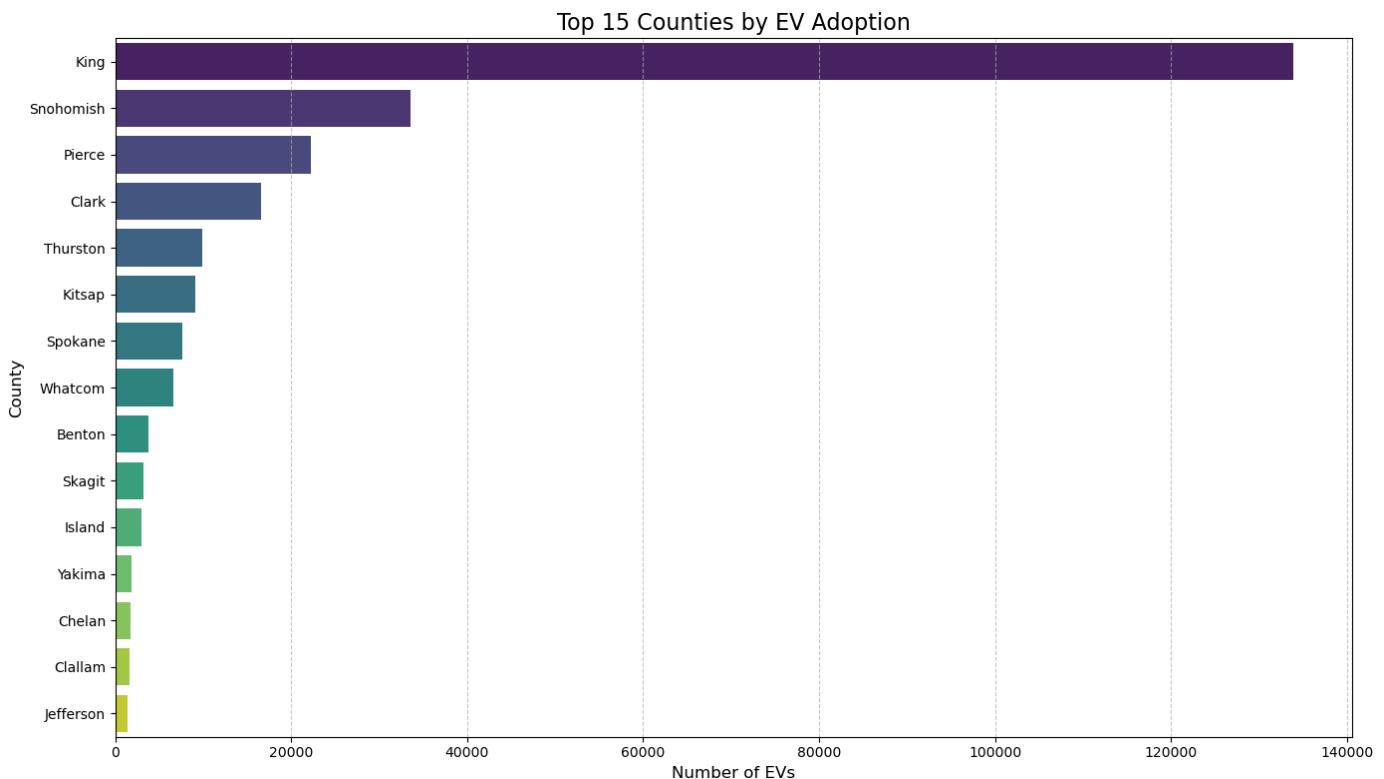
C:\Users\LENOVO\AppData\Local\Temp\ipykernel_1800\3802896159.py:22: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0.
Assign the `y` variable to `hue` and set `legend=False` for the same effect.

```

sns.barplot(x='Number of EVs', y='County', data=ev_by_county.head(15), palette='viridis')

```



DATA VISUALISATION

Create a bar chart showing the top 5 EV makes and models by count.

In [34]:

```

import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd

# Count EVs by make and model
make_model_counts = df1.groupby(['Make', 'Model']).size().reset_index(name='Count')

# Sort by count in descending order
make_model_counts = make_model_counts.sort_values('Count', ascending=False)

# Get the top 5 make-model combinations
top_5_make_models = make_model_counts.head(5)

# Create a new column combining make and model for display
top_5_make_models['Make_Model'] = top_5_make_models['Make'] + ' ' + top_5_make_models['Model']

# Create the bar chart
plt.figure(figsize=(12, 6))
bars = sns.barplot(x='Make_Model', y='Count', data=top_5_make_models, palette='viridis')

# Add title and labels
plt.title('Top 5 EV Makes and Models by Count', fontsize=16)
plt.xlabel('Make and Model', fontsize=12)
plt.ylabel('Number of Vehicles', fontsize=12)
plt.xticks(rotation=45, ha='right')

# Add value labels on top of each bar
for i, v in enumerate(top_5_make_models['Count']):
    plt.text(i, v + 500, str(v), rotation=90, va='bottom')

```

```

    plt.text(i, v + 0.5, str(v), ha='center')

# Add grid lines for better readability
plt.grid(axis='y', linestyle='--', alpha=0.7)

# Adjust layout
plt.tight_layout()

# Show the plot
plt.show()

# Print the data for reference
print("Top 5 EV Makes and Models by Count:")
print(top_5_make_models[['Make', 'Model', 'Count']])

# Calculate percentage of total
total_evs = len(df1)
top_5_make_models['Percentage'] = (top_5_make_models['Count'] / total_evs * 100).round(2)
print("\nPercentage of Total EVs:")
print(top_5_make_models[['Make_Model', 'Count', 'Percentage']])

```

C:\Users\LENOVO\AppData\Local\Temp\ipykernel_1800\2335216976.py:15: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row_indexer,col_indexer] = value instead

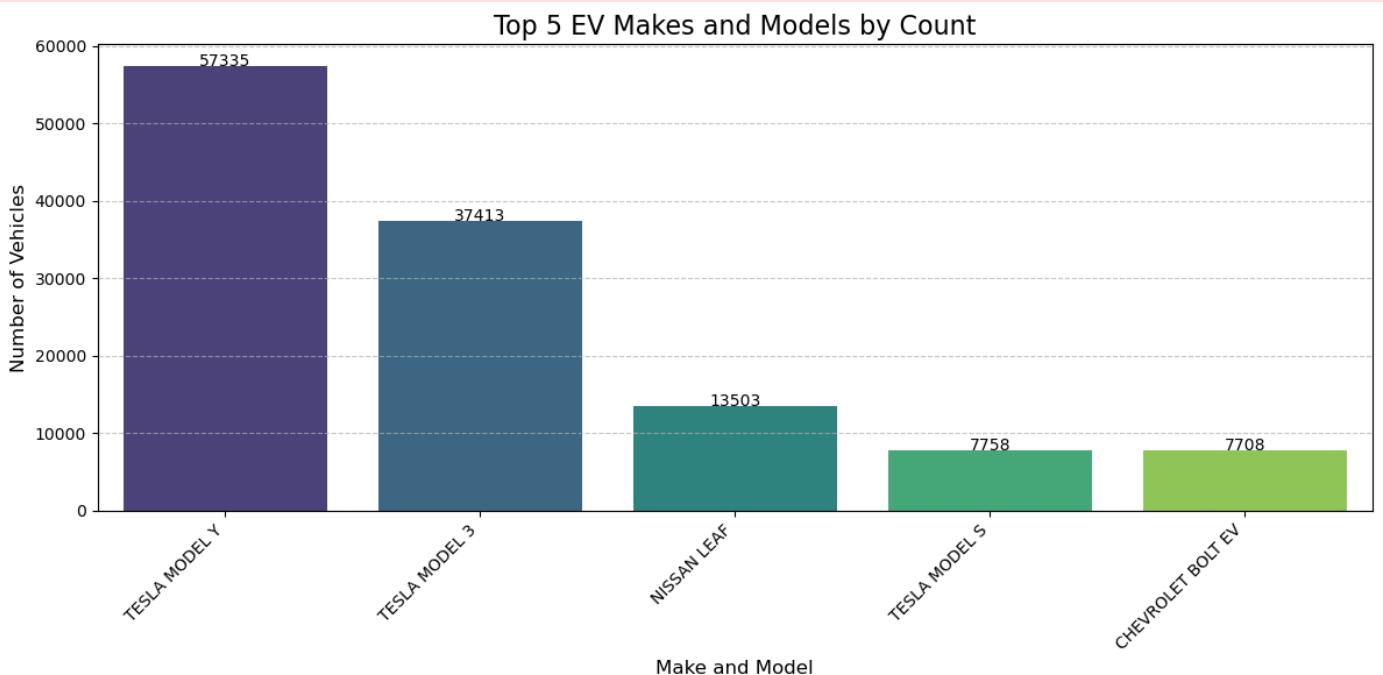
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
top_5_make_models['Make_Model'] = top_5_make_models['Make'] + ' ' + top_5_make_models['Model']
```

C:\Users\LENOVO\AppData\Local\Temp\ipykernel_1800\2335216976.py:19: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
bars = sns.barplot(x='Make_Model', y='Count', data=top_5_make_models, palette='viridis')
```



Top 5 EV Makes and Models by Count:
 Make Model Count

```
159      TESLA MODEL Y  57335
156      TESLA MODEL 3  37413
136      NISSAN LEAF   13503
157      TESLA MODEL S  7758
48      CHEVROLET BOLT EV  7708
```

Percentage of Total EVs:

	Make_Model	Count	Percentage
159	TESLA MODEL Y	57335	21.21
156	TESLA MODEL 3	37413	13.84
136	NISSAN LEAF	13503	5.00
157	TESLA MODEL S	7758	2.87
48	CHEVROLET BOLT EV	7708	2.85

```
C:\Users\LENOVO\AppData\Local\Temp\ipykernel_1800\2335216976.py:46: SettingWithCopyWarning:
```

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
top_5_make_models['Percentage'] = (top_5_make_models['Count'] / total_evs * 100).round(2)
```

Use a heatmap or choropleth map to visualize EV distribution by county.

In [35]:

```
# pip install geopandas
```

In [36]:

```
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd
import numpy as np
import geopandas as gpd
from matplotlib.colors import LinearSegmentedColormap

# Step 1: Count EVs by county
if 'County' in df1.columns:
    # Count EVs by county
    ev_by_county = df1['County'].value_counts().reset_index()
    ev_by_county.columns = ['County', 'Number of EVs']

    print("EV Distribution by County:")
    print(ev_by_county.head(10))

# Step 2: Create a heatmap-style visualization
# First, ensure county names are standardized (e.g., remove "County" suffix if present)
ev_by_county['County'] = ev_by_county['County'].str.replace(' County', '', regex=False)

# Create a pivot table for the heatmap
# We'll use a dummy column since we only have one value per county
ev_by_county['Dummy'] = 'EV Count'
heatmap_data = ev_by_county.pivot(index='County', columns='Dummy', values='Number of EVs')

# Sort by EV count for better visualization
heatmap_data = heatmap_data.sort_values('EV Count', ascending=False)

# Create a heatmap
```

```

plt.figure(figsize=(10, max(8, len(ev_by_county) * 0.3))) # Adjust height based on
sns.heatmap(heatmap_data, annot=True, fmt='d', cmap='YlGnBu',
            linewidths=.5, cbar_kws={'label': 'Number of EVs'})
plt.title('EV Distribution by County (Heatmap)', fontsize=16)
plt.tight_layout()
plt.show()

# Step 3: Create a choropleth map
try:
    # Load US counties shapefile
    # You may need to download this file or use a different source
    # This example uses the US Census Bureau's county shapefile
    counties = gpd.read_file('https://raw.githubusercontent.com/plotly/datasets/master/us-counties.shp')

    # If the above URL doesn't work, you can use:
    # counties = gpd.read_file(gpd.datasets.get_path('us_counties'))

    # Alternatively, for a specific state, you might use:
    # counties = gpd.read_file('path_to_state_county_shapefile.shp')

    # Standardize county names for joining
    counties['NAME'] = counties['NAME'].str.upper()
    ev_by_county['County'] = ev_by_county['County'].str.upper()

    # Merge EV data with geographic data
    # Note: The actual column names in your shapefile may differ
    merged_data = counties.merge(ev_by_county, left_on='NAME', right_on='County', how='left')

    # Fill NaN values with 0 (counties with no EVs in the dataset)
    merged_data['Number of EVs'] = merged_data['Number of EVs'].fillna(0)

    # Create the choropleth map
    fig, ax = plt.subplots(1, 1, figsize=(15, 10))

    # Create a custom colormap
    cmap = LinearSegmentedColormap.from_list("", ["#f7fbff", "#08519c"])

    # Plot the choropleth
    merged_data.plot(column='Number of EVs',
                      ax=ax,
                      legend=True,
                      cmap=cmap,
                      legend_kwds={'label': "Number of EVs by County",
                                   'orientation': "horizontal"})

    # Add title and remove axis
    ax.set_title('EV Distribution by County', fontsize=16)
    ax.set_axis_off()

    plt.tight_layout()
    plt.show()

except Exception as e:
    print(f"Could not create choropleth map: {e}")
    print("To create a choropleth map, you need to install geopandas and have access")
    print("Alternative approach: Consider using Folium or Plotly for interactive map")

# Alternative: Create a simple bar chart of top counties
plt.figure(figsize=(14, 8))

```

```

top_counties = ev_by_county.head(15)
sns.barplot(x='Number of EVs', y='County', data=top_counties, palette='YlGnBu')
plt.title('Top 15 Counties by EV Count', fontsize=16)
plt.xlabel('Number of EVs', fontsize=12)
plt.ylabel('County', fontsize=12)
plt.grid(axis='x', linestyle='--', alpha=0.7)
plt.tight_layout()
plt.show()

# Provide code for interactive map using Plotly
print("\nHere's code for an interactive choropleth map using Plotly:")
print("""
import plotly.express as px

# Load county FIPS codes (you'll need to create or obtain this mapping)
# county_fips = pd.read_csv('county_fips.csv') # Map county names to FIPS codes
# ev_with_fips = pd.merge(ev_by_county, county_fips, on='County', how='left')

# For this example, we'll use a sample dataset with FIPS codes
fig = px.choropleth(ev_by_county,
                     geojson=counties,
                     locations='FIPS', # FIPS code column
                     color='Number of EVs',
                     color_continuous_scale="Viridis",
                     scope="usa",
                     labels={'Number of EVs': 'EV Count'}
                    )
fig.update_layout(margin={"r":0,"t":0,"l":0,"b":0})
fig.show()
""")

else:
    print("No 'County' column found in the dataset.")
    print("Available columns:", df1.columns.tolist())
    print("\nTo create a county-level visualization, you need county data in your database")
    print("If you have ZIP codes or other location data, you could aggregate to county level")

```

EV Distribution by County:

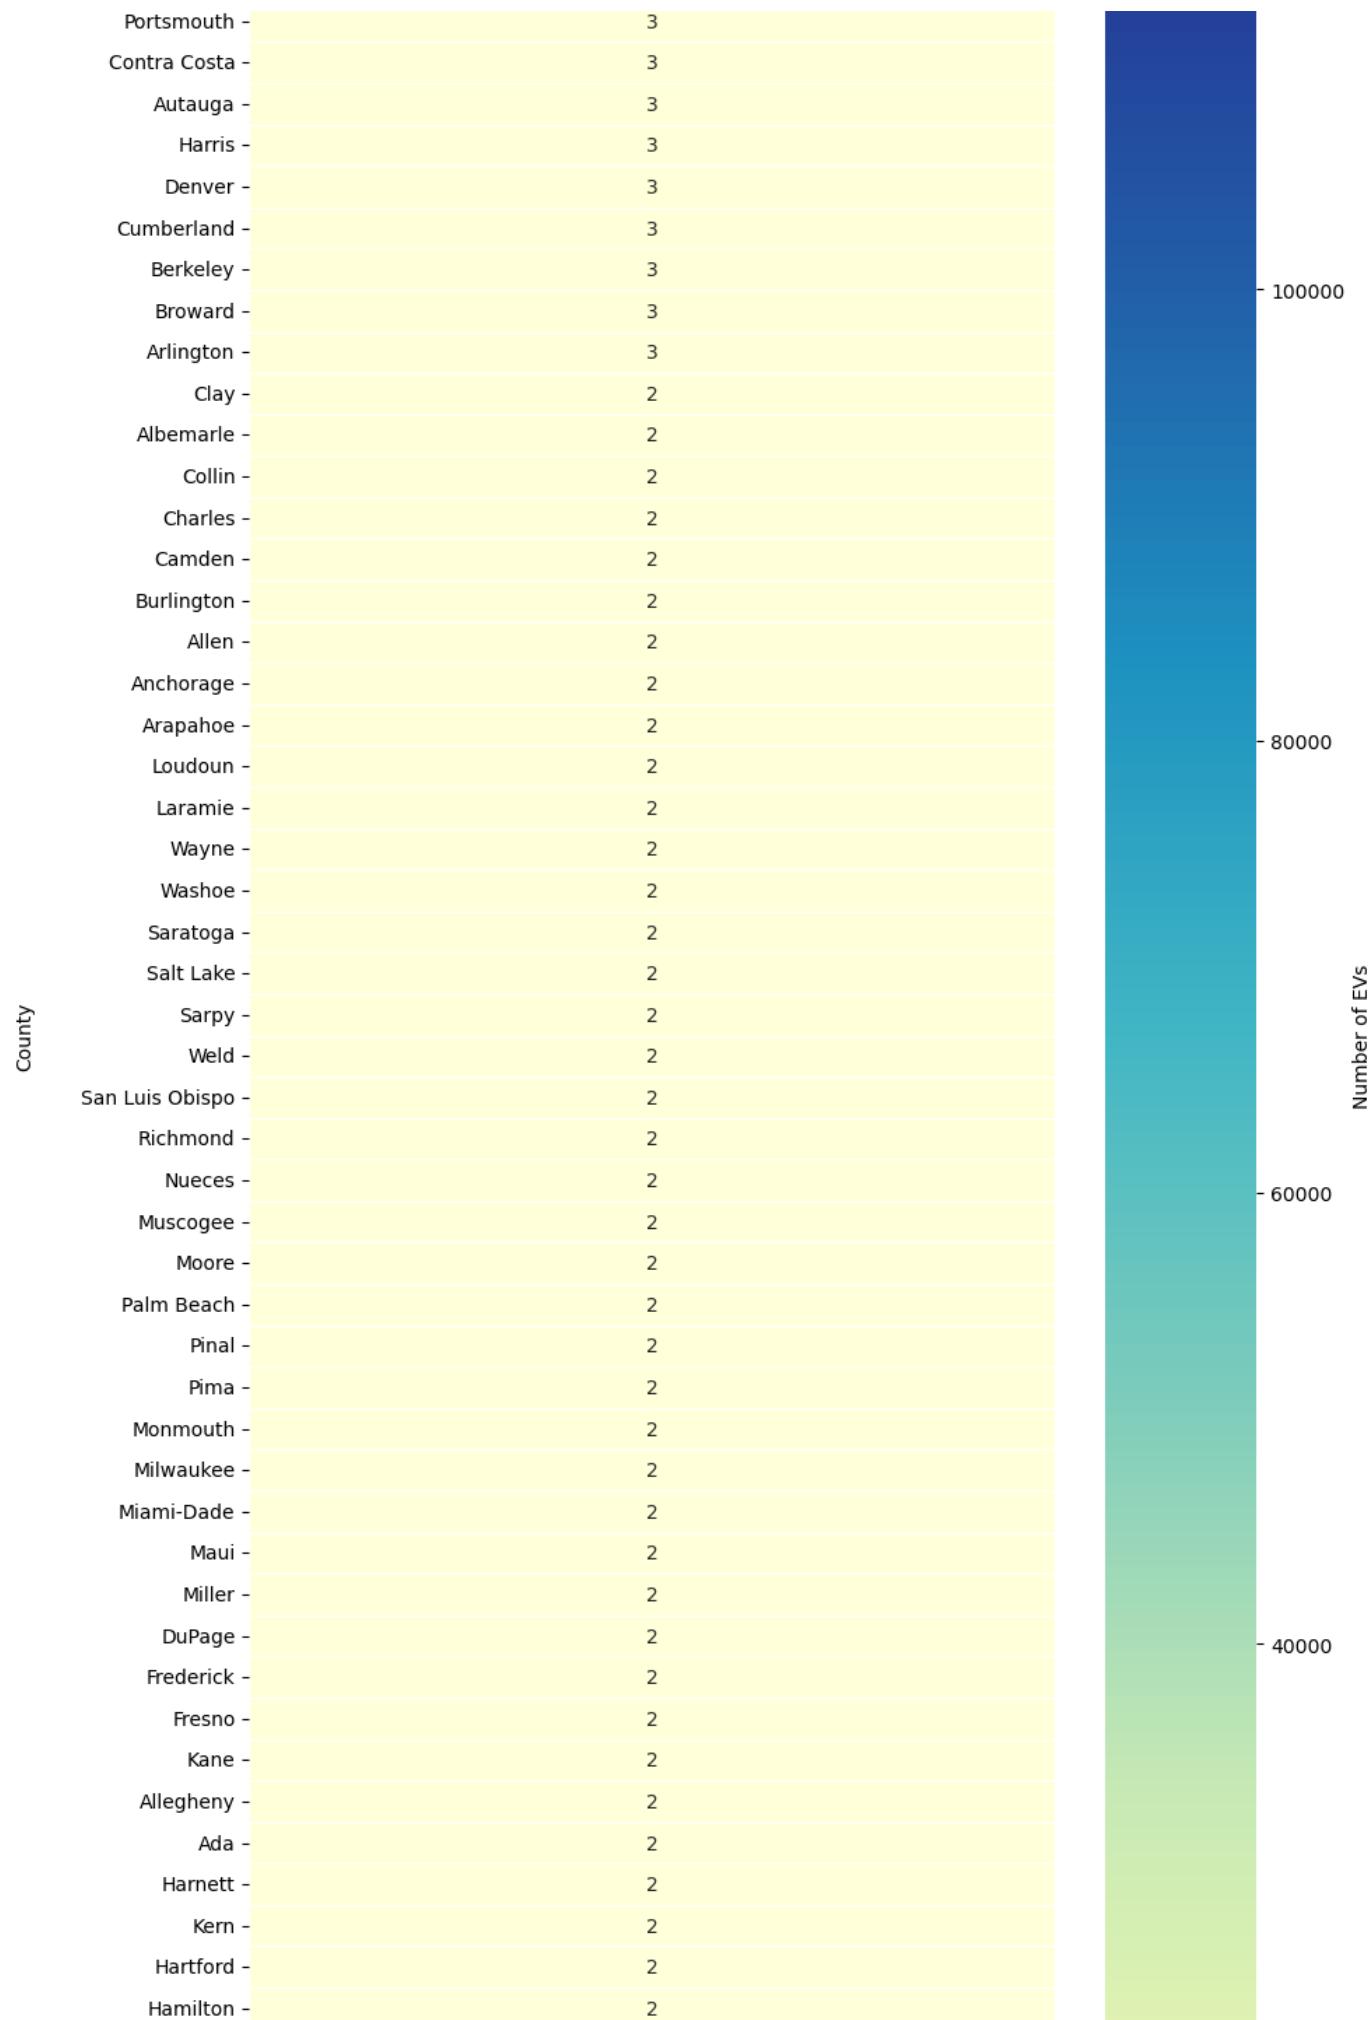
	County	Number of EVs
0	King	133903
1	Snohomish	33531
2	Pierce	22213
3	Clark	16553
4	Thurston	9852
5	Kitsap	9057
6	Spokane	7593
7	Whatcom	6620
8	Benton	3792
9	Skagit	3166

EV Distribution by County (Heatmap)

King -	133903
Snohomish -	33531
Pierce -	22213
Clark -	16553
Thurston -	9852
Kitsap -	9057
Spokane -	7593
Whatcom -	6620
Benton -	3792
Skagit -	3166
Island -	2962
Yakima -	1864
Chelan -	1691
Clallam -	1647
Jefferson -	1435
Cowlitz -	1401
Mason -	1359
San Juan -	1239
Lewis -	1221
Franklin -	1155
Grant -	1053
Grays Harbor -	1034
Kittitas -	1017
Walla Walla -	733
Douglas -	610
Whitman -	542
Klickitat -	474
Okanogan -	395
Stevens -	347
Pacific -	340
Skamania -	265
Asotin -	114
Adams -	111
Wahkiakum -	95
Pend Oreille -	88
Lincoln -	79
Ferry -	50
San Diego -	30
Columbia -	28
Fairfax -	24
El Paso -	22
Orange -	21
Santa Clara -	16
Los Angeles -	15
Riverside -	14
Anne Arundel -	12
Multnomah -	12
Maricopa -	11

Marietta -	11
Alameda -	11
Duval -	8
Bexar -	8
Lake -	8
Montgomery -	8
Virginia Beach -	7
Kings -	7
Alexandria -	7
Macomb -	6
Garfield -	6
Chesapeake -	6
Leavenworth -	6
Monterey -	6
Prince George's -	6
Hillsborough -	5
District of Columbia -	5
Clackamas -	5
Harford -	5
San Bernardino -	5
San Mateo -	5
Rockingham -	5
New London -	5
Middlesex -	5
Prince William -	5
York -	4
Williamson -	4
Travis -	4
Norfolk -	4
Okaloosa -	4
Ventura -	4
Stafford -	4
Cook -	4
New York -	4
Kootenai -	4
Bell -	4
St. Louis -	3
Washington -	3
Newport -	3
Polk -	3
Placer -	3
Marin -	3
Pulaski -	3
San Francisco -	3
Shelby -	3
Suffolk -	3
St. Mary's -	3
Solano -	3
St. Clair -	3

- 120000

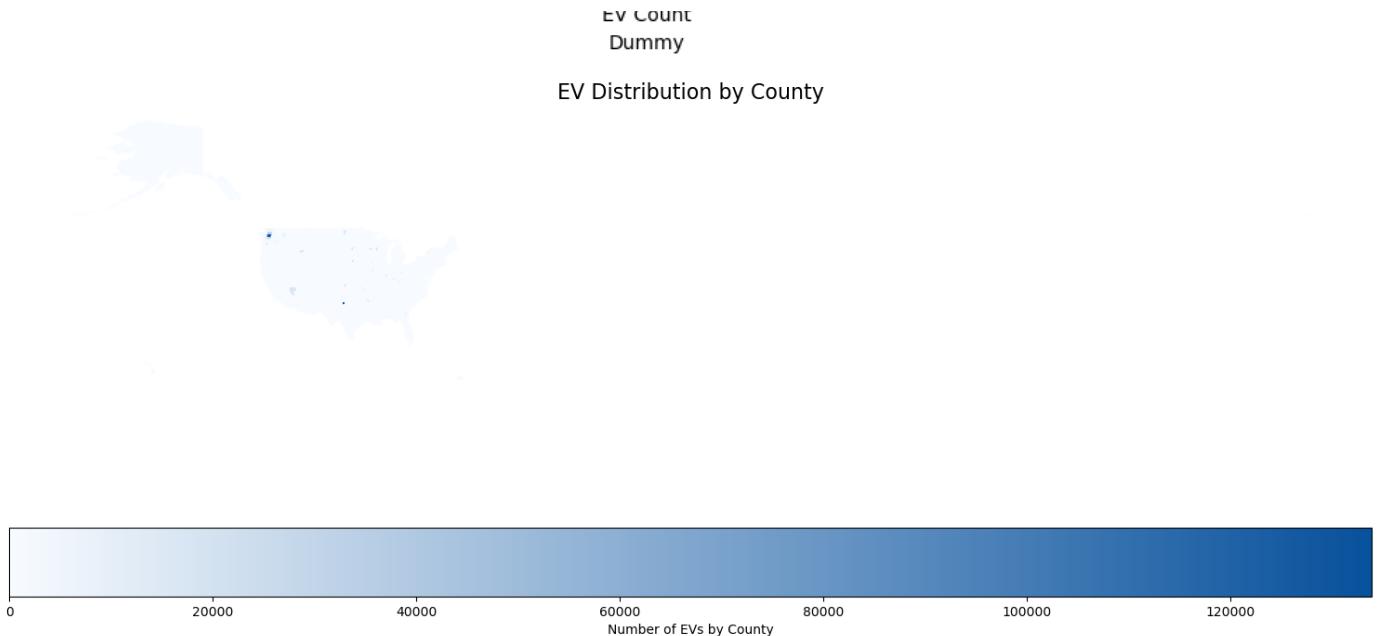




Kent -	2
Dallas -	2
DeKalb -	2
Howard -	2
Hudson -	2
Essex -	2
Beauregard -	2
Brevard -	2
Bucks -	2
Calvert -	1
Calhoun -	1
Carroll -	1
Baltimore -	1
Atlantic -	1
Beaver -	1
Boulder -	1
Brown -	1
Beaufort -	1
Barnstable -	1
Bannock -	1
Caddo -	1
Bernalillo -	1
Currituck -	1
Cuyahoga -	1
Dale -	1
Escambia -	1
Fort Bend -	1
Forsyth -	1
Centre -	1
Churchill -	1
Chesterfield -	1
Cochise -	1
Cobb -	1
Johnson -	1
Honolulu -	1
Jackson -	1
Hoke -	1
Houston -	1
James City -	1
Hennepin -	1
Greenville -	1
Hardin -	1
Hampden -	1
Gwinnett -	1
Guam -	1
Greene -	1
Harrison -	1
Henrico -	1
Galveston -	1

Deschutes -	1
Denton -	1
Davidson -	1
Doña Ana -	1
Coryell -	1
Larimer -	1
Lexington -	1
Lee -	1
Manassas -	1
Madison -	1
Manatee -	1
Marion -	1
Medina -	1
Meade -	1
Mercer -	1
Monroe -	1
Prince George -	1
Platte -	1
Otero -	1
Penobscot -	1
Osceola -	1
Niagara -	1
Ocean -	1
Oldham -	1
Nye -	1
Nassau -	1
Saginaw -	1
Sacramento -	1
Sarasota -	1
Santa Cruz -	1
Riley -	1
San Joaquin -	1
St. Landry -	1
St. Johns -	1
St. Charles -	1
Sonoma -	1
Talladega -	1
Stanislaus -	1
Tarrant -	1
Tooele -	1
Umatilla -	1
Texas -	1
Tom Green -	1
Wake -	1
Weber -	1
Washtenaw -	1
Worcester -	1
Yuba -	1

Page 1



Create a line graph showing the trend of EV adoption by model year.

In [37]:

```
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd
import numpy as np

# Count EVs by model year
ev_by_year = df1['Model Year'].value_counts().sort_index()

# Convert to DataFrame for easier manipulation
ev_year_df = ev_by_year.reset_index()
ev_year_df.columns = ['Model Year', 'Number of EVs']

# Create the line graph
plt.figure(figsize=(14, 8))

# Plot the line with markers
plt.plot(ev_year_df['Model Year'], ev_year_df['Number of EVs'],
         marker='o', linestyle='--', linewidth=2, markersize=8, color="#1f77b4")

# Add data points
for x, y in zip(ev_year_df['Model Year'], ev_year_df['Number of EVs']):
    plt.text(x, y + max(ev_year_df['Number of EVs'])*0.02, f'{y:,}', ha='center', va='bottom', fontsize=9)

# Add title and labels
plt.title('EV Adoption Trend by Model Year', fontsize=16)
plt.xlabel('Model Year', fontsize=12)
plt.ylabel('Number of EVs', fontsize=12)

# Add grid for better readability
plt.grid(True, linestyle='--', alpha=0.7)

# Set x-axis ticks to show all years
plt.xticks(ev_year_df['Model Year'], rotation=45)

# Calculate year-over-year growth rates
```

```

ev_year_df['YoY Growth'] = ev_year_df['Number of EVs'].pct_change() * 100

# Add a second y-axis for growth rate
ax2 = plt.gca().twinx()
ax2.plot(ev_year_df['Model Year'][1:], ev_year_df['YoY Growth'][1:],
         marker='s', linestyle='--', color='green', alpha=0.7)
ax2.set_ylabel('Year-over-Year Growth (%)', color='green', fontsize=12)
ax2.tick_params(axis='y', labelcolor='green')

# Add legend
plt.legend(['Number of EVs', 'YoY Growth (%)'], loc='upper left')

# Adjust layout
plt.tight_layout()

# Show the plot
plt.show()

# Print the data for reference
print("EV Adoption by Model Year:")
print(ev_year_df)

# Calculate CAGR (Compound Annual Growth Rate)
if len(ev_year_df) > 1:
    first_year = ev_year_df['Model Year'].min()
    last_year = ev_year_df['Model Year'].max()
    first_year_count = ev_year_df[ev_year_df['Model Year'] == first_year]['Number of EVs']
    last_year_count = ev_year_df[ev_year_df['Model Year'] == last_year]['Number of EVs']

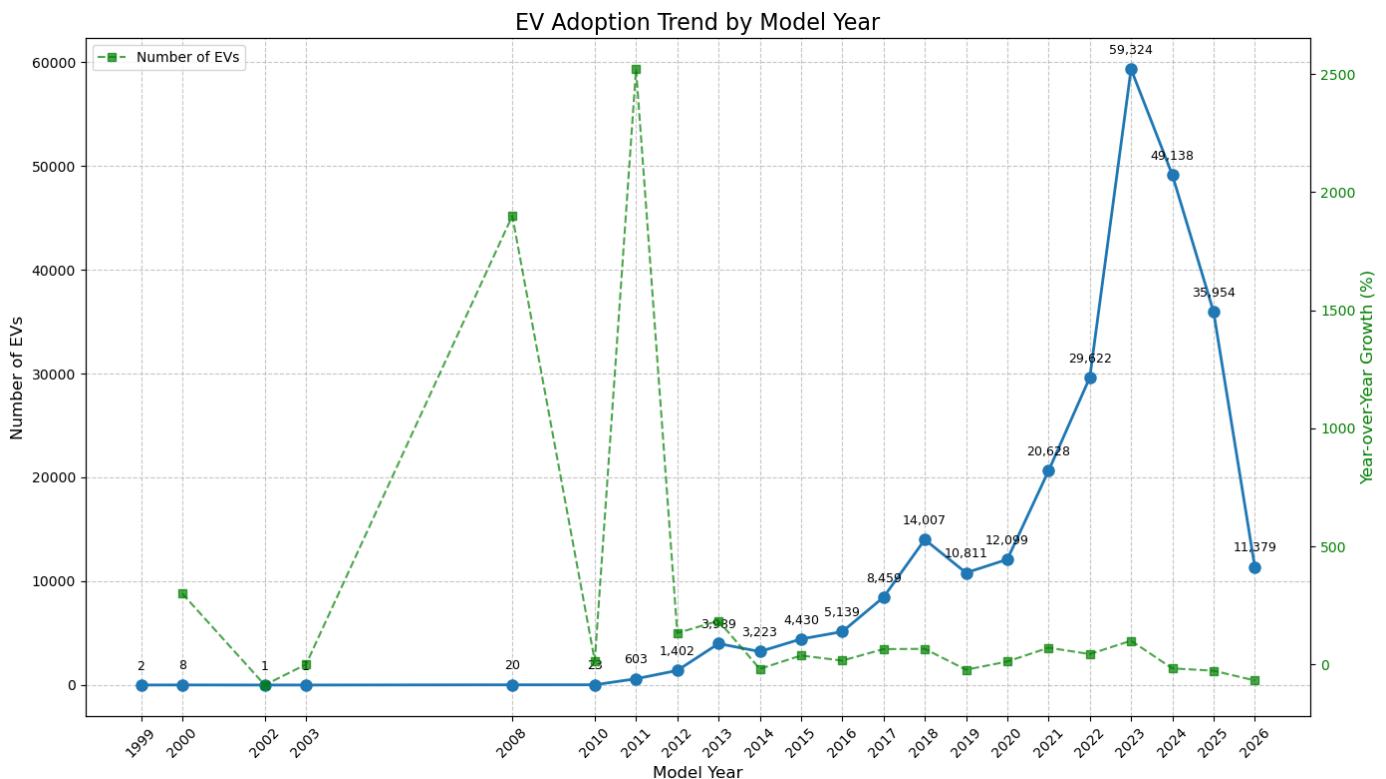
    # Only calculate if first year count is greater than 0
    if first_year_count > 0:
        years = last_year - first_year
        cagr = (((last_year_count / first_year_count) ** (1 / years)) - 1) * 100
        print(f"\nCompound Annual Growth Rate (CAGR) from {first_year} to {last_year}: {cagr:.2f}%")

# Create a bar chart version for comparison
plt.figure(figsize=(14, 8))
sns.barplot(x='Model Year', y='Number of EVs', data=ev_year_df, palette='viridis')

# Add value labels on top of each bar
for i, v in enumerate(ev_year_df['Number of EVs']):
    plt.text(i, v + max(ev_year_df['Number of EVs'])*0.02, f'{v:,}', ha='center')

plt.title('EV Adoption by Model Year (Bar Chart)', fontsize=16)
plt.xlabel('Model Year', fontsize=12)
plt.ylabel('Number of EVs', fontsize=12)
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()

```



EV Adoption by Model Year:

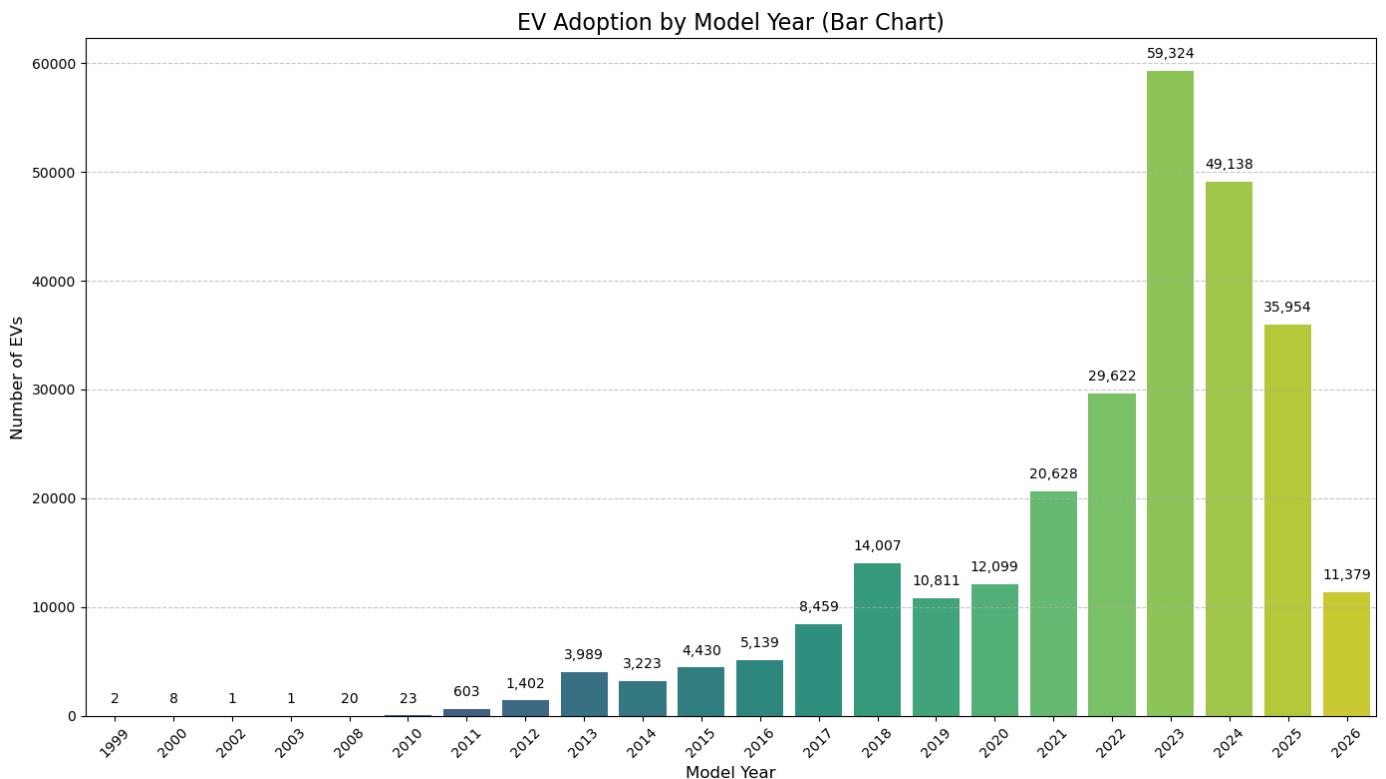
	Model Year	Number of EVs	YoY Growth
0	1999	2	Nan
1	2000	8	300.000000
2	2002	1	-87.500000
3	2003	1	0.000000
4	2008	20	1900.000000
5	2010	23	15.000000
6	2011	603	2521.739130
7	2012	1402	132.504146
8	2013	3989	184.522111
9	2014	3223	-19.202808
10	2015	4430	37.449581
11	2016	5139	16.004515
12	2017	8459	64.604009
13	2018	14007	65.586949
14	2019	10811	-22.817163
15	2020	12099	11.913792
16	2021	20628	70.493429
17	2022	29622	43.600931
18	2023	59324	100.270070
19	2024	49138	-17.170117
20	2025	35954	-26.830559
21	2026	11379	-68.351227

Compound Annual Growth Rate (CAGR) from 1999 to 2026: 37.75%

C:\Users\LENOVO\AppData\Local\Temp\ipykernel_1800\1767253685.py:74: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(x='Model Year', y='Number of EVs', data=ev_year_df, palette='viridis')
```



Plot a pie chart showing the proportion of CAFV-eligible vs. non-eligible EVs.

In [38]:

```
import matplotlib.pyplot as plt
import pandas as pd
import numpy as np

# Check if the CAFV eligibility column exists
# Common column names for this information
possible_columns = [
    'Clean Alternative Fuel Vehicle Eligible',
    'Eligibility unknown as battery range has not been researched']

# Find the correct column name
cafv_column = None
for col in possible_columns:
    if col in df1.columns:
        cafv_column = col
        break

if cafv_column:
    # Check the unique values in the column
    print(f"Unique values in {cafv_column}:")
    print(df1[cafv_column].value_counts())

    # Count the number of eligible and non-eligible vehicles
    cafv_counts = df1[cafv_column].value_counts()

    # Determine which values indicate eligibility
    # This depends on how the data is coded in your dataset
    # Common values might be 'Eligible', 'Yes', 'Clean Alternative Fuel Vehicle', etc.

    # For this example, we'll assume 'Eligible' indicates eligibility
    # Adjust this based on your actual data values
```

```

eligible_values = ['Clean Alternative Fuel Vehicle Eligible',
    'Eligibility unknown as battery range has not been researched']

# Create a simplified DataFrame for the pie chart
# Group all eligible values together and all non-eligible values together
eligibility_data = pd.DataFrame({
    'Status': ['Eligible', 'Not Eligible'],
    'Count': [
        df1[cafv_column].isin(eligible_values).sum(),
        (~df1[cafv_column].isin(eligible_values)).sum()
    ]
})

# Calculate percentages
eligibility_data['Percentage'] = (eligibility_data['Count'] / eligibility_data['Count']).round(2)

# Create labels with percentages
labels = [f"{status} ({count:,}), {pct}%"
          for status, count, pct in zip(eligibility_data['Status'],
                                         eligibility_data['Count'],
                                         eligibility_data['Percentage'])]

# Create the pie chart
plt.figure(figsize=(10, 8))

# Plot the pie chart with a slight explosion for the eligible segment
explode = (0.1, 0) # explode the 1st slice (Eligible)

plt.pie(eligibility_data['Count'],
        explode=explode,
        labels=labels,
        autopct=' ', # We're using custom labels with percentages
        startangle=90,
        shadow=True,
        colors=['#66b3ff', '#ff9999'],
        wedgeprops={'edgecolor': 'white', 'linewidth': 1.5})

# Equal aspect ratio ensures that pie is drawn as a circle
plt.axis('equal')

# Add title
plt.title('Proportion of CAFV-Eligible vs. Non-Eligible EVs', fontsize=16)

# Add a legend
plt.legend(eligibility_data['Status'], loc="best")

# Show the plot
plt.tight_layout()
plt.show()

# Print the summary data
print("\nCAFV Eligibility Summary:")
print(eligibility_data)

# Additional analysis: Check if eligibility varies by make
print("\nCAFV Eligibility by Make:")
make_eligibility = df1.groupby('Make')[cafv_column].apply(
    lambda x: (x.isin(eligible_values).sum() / len(x) * 100).round(1)
).sort_values(ascending=False)

```

```

print(make_eligibility)

# Create a bar chart showing eligibility percentage by make
plt.figure(figsize=(14, 8))
make_eligibility.plot(kind='bar', color='skyblue')
plt.title('Percentage of CAFV-Eligible EVs by Make', fontsize=16)
plt.xlabel('Make', fontsize=12)
plt.ylabel('Percentage Eligible (%)', fontsize=12)
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.xticks(rotation=45, ha='right')

# Add value labels on top of each bar
for i, v in enumerate(make_eligibility):
    plt.text(i, v + 1, f"{v}%", ha='center')

plt.tight_layout()
plt.show()

else:
    print("No CAFV eligibility column found in the dataset.")
    print("Available columns:", df1.columns.tolist())

```

No CAFV eligibility column found in the dataset.

Available columns: ['VIN (1-10)', 'County', 'City', 'State', 'Postal Code', 'Model Year', 'Make', 'Model', 'Electric Vehicle Type', 'Clean Alternative Fuel Vehicle (CAFV) Eligibility', 'Electric Range', 'Legislative District', 'DOL Vehicle ID', 'Vehicle Location', 'Electric Utility', '2020 Census Tract']

In [39]:

```
!pip install folium
```

```
Requirement already satisfied: folium in c:\users\lenovo\anaconda3\lib\site-packages (0.20.0)
Requirement already satisfied: branca>=0.6.0 in c:\users\lenovo\anaconda3\lib\site-packages (from folium) (0.8.2)
Requirement already satisfied: jinja2>=2.9 in c:\users\lenovo\anaconda3\lib\site-packages (from folium) (3.1.6)
Requirement already satisfied: numpy in c:\users\lenovo\anaconda3\lib\site-packages (from folium) (2.1.3)
Requirement already satisfied: requests in c:\users\lenovo\anaconda3\lib\site-packages (from folium) (2.32.3)
Requirement already satisfied: xyzservices in c:\users\lenovo\anaconda3\lib\site-packages (from folium) (2022.9.0)
Requirement already satisfied: MarkupSafe>=2.0 in c:\users\lenovo\anaconda3\lib\site-packages (from jinja2>=2.9->folium) (3.0.2)
Requirement already satisfied: charset-normalizer<4,>=2 in c:\users\lenovo\anaconda3\lib\site-packages (from requests->folium) (3.3.2)
Requirement already satisfied: idna<4,>=2.5 in c:\users\lenovo\anaconda3\lib\site-packages (from requests->folium) (3.7)
Requirement already satisfied: urllib3<3,>=1.21.1 in c:\users\lenovo\anaconda3\lib\site-packages (from requests->folium) (2.3.0)
Requirement already satisfied: certifi>=2017.4.17 in c:\users\lenovo\anaconda3\lib\site-packages (from requests->folium) (2025.8.3)
```

Use a geospatial map to display EV registrations based on vehicle location.

In []:

```

# Import the necessary libraries first
import pandas as pd
import plotly.express as px # This import was missing
import folium
from folium.plugins import MarkerCluster # This import might also be needed

# Rest of your code remains the same
# We're using the right ones
print(df1.columns.tolist())

# We need to extract latitude and longitude from 'Vehicle Location' column
# Assuming 'Vehicle Location' contains coordinates in format "POINT (longitude latitude)"
# Extract coordinates from 'Vehicle Location'
coords = df1['Vehicle Location'].str.extract(r'POINT \((([-\d.]+) ([-\d.]+)\))')
# Rename the extracted columns
coords.columns = ['Longitude', 'Latitude']
# Convert to numeric and assign to df1
df1['Longitude'] = pd.to_numeric(coords['Longitude'])
df1['Latitude'] = pd.to_numeric(coords['Latitude'])

# Method 1: Using Plotly Express (interactive and works well in Jupyter)
fig = px.scatter_mapbox(df1,
                        lat='Latitude',
                        lon='Longitude',
                        hover_name='City',
                        hover_data=['Make', 'Model'],
                        color='Make',
                        zoom=8,
                        height=600,
                        width=800,
                        title='EV Registrations by Location')

fig.update_layout(mapbox_style="open-street-map")
fig.update_layout(margin={"r":0,"t":50,"l":0,"b":0})
fig.show()

# Method 2: Using Folium (more customizable)
# Filter out rows with NaN values in Latitude or Longitude
# This is the key fix to address the error
df_clean = df1.dropna(subset=['Latitude', 'Longitude'])

# Create a map centered at the mean of your data points
m = folium.Map(location=[df_clean['Latitude'].mean(), df_clean['Longitude'].mean()],
               zoom_start=10)

# Add a marker cluster to make the map more readable with many points
marker_cluster = MarkerCluster().add_to(m)

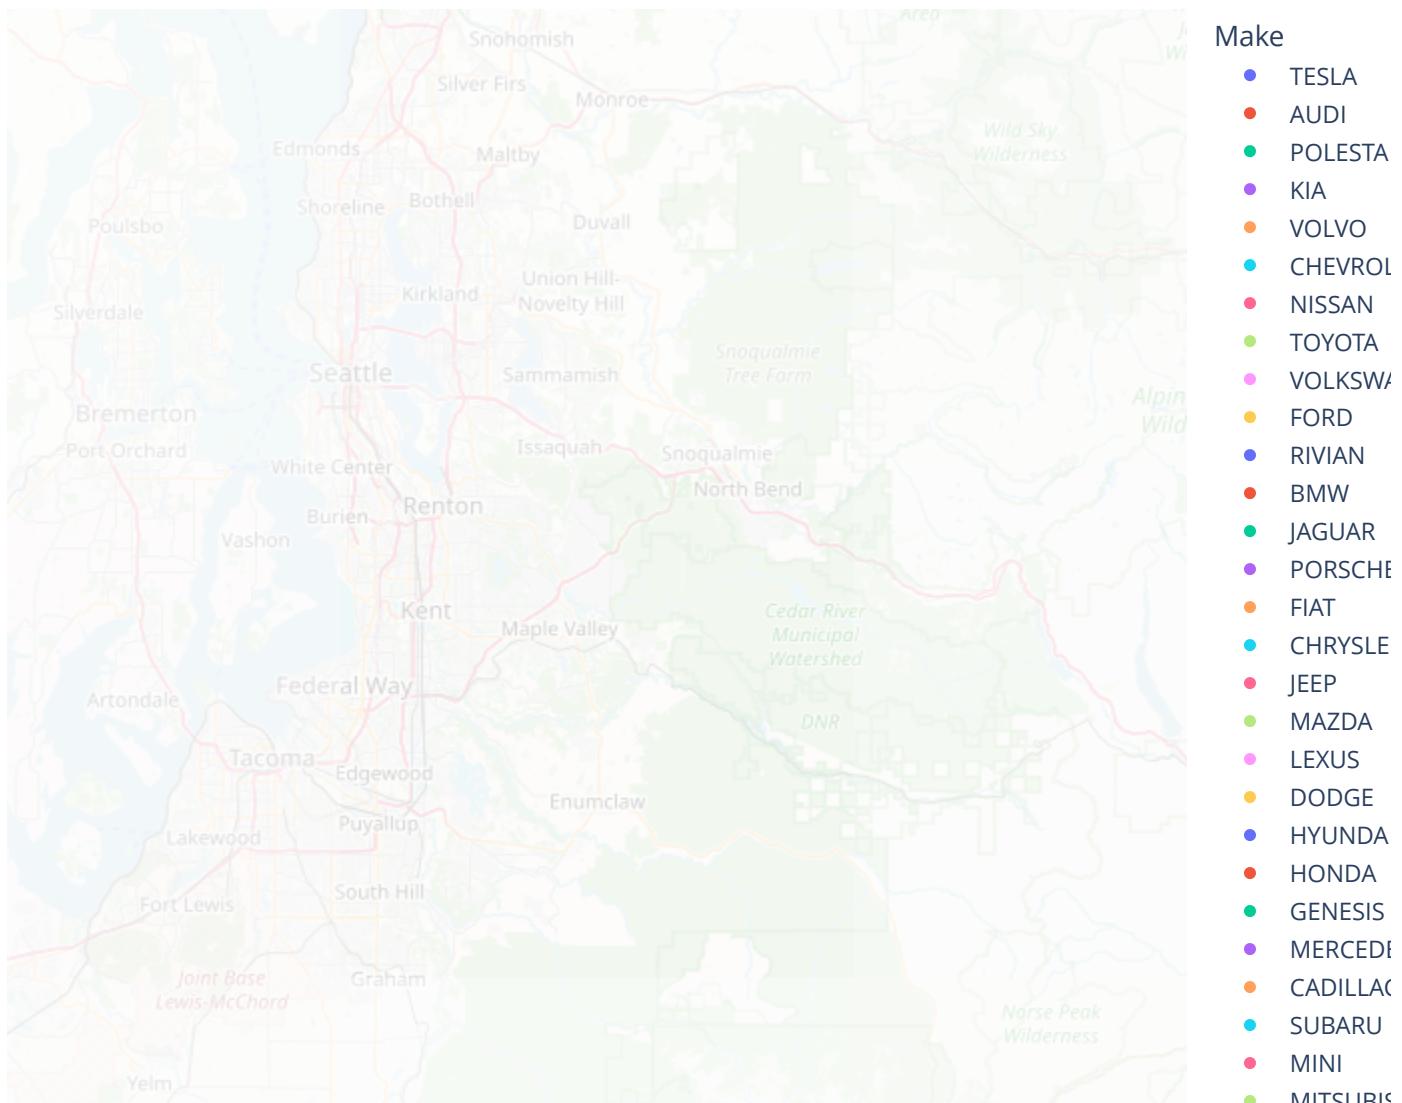
# Add markers for each EV registration
for idx, row in df_clean.iterrows(): # Using the filtered dataframe
    folium.Marker(
        location=[row['Latitude'], row['Longitude']],
        popup=f"Make: {row['Make']}<br>Model: {row['Model']}<br>City: {row['City']}",
        tooltip=row['City']
    ).add_to(marker_cluster)

# Display the map
m

```

```
['VIN (1-10)', 'County', 'City', 'State', 'Postal Code', 'Model Year', 'Make', 'Model',  
'Electric Vehicle Type', 'Clean Alternative Fuel Vehicle (CAFV) Eligibility', 'Electric  
Range', 'Legislative District', 'DOL Vehicle ID', 'Vehicle Location', 'Electric Utilit  
y', '2020 Census Tract']
```

EV Registrations by Location



In []:

```
# pip install scikit-learn
```

What independent variables (features) can be used to predict Electric Range? (e.g., Model Year, Base MSRP, Make)

In []:

Linear Regression

```
In [27]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score

In [6]: df1 = pd.read_csv("D:\manish\Python\GROW AI\python\Electric_Vehicle_Population_Data.csv")

Out[6]:
```

VIN (1-10)	County	City	State	Postal Code	Model Year	Make	Model	Electric Vehicle Type	Clean Alternative Fuel Vehicle (CAVF) Eligibility	Electric Range	Legislative District	DOL	Vehicle ID	Vehicle Location	Electric Utility	2020 Census Tract
0 5YJYGDDE8L	Thurston	Tumwater	WA	98501.0	2020	TESLA	MODEL Y	Battery Electric Vehicle (BEV)	Clean Alternative Fuel Vehicle Eligible	291.0	35.0	124633715	POINT	PUGET SOUND ENERGY INC	5.306701e+10	
1 5YXCAE2XJ	Snohomish	Bothell	WA	98021.0	2018	TESLA	MODEL X	Battery Electric Vehicle (BEV)	Clean Alternative Fuel Vehicle Eligible	238.0	1.0	474826075	POINT	PUGET SOUND ENERGY INC	5.306105e+10	
2 5YJ3E1EBXK	King	Kent	WA	98031.0	2019	TESLA	MODEL 3	Battery Electric Vehicle (BEV)	Clean Alternative Fuel Vehicle Eligible	220.0	47.0	280302733	POINT	PUGET SOUND ENERGY INC	5.303303e+10	
3 7SAYGDEEAT	King	Issaquah	WA	98027.0	2026	TESLA	MODEL Y	Battery Electric Vehicle (BEV)	Eligibility unknown as battery range has not b...	0.0	41.0	280786565	POINT	PUGET SOUND ENERGY INC	5.303302e+10	
4 WAUUPBFF9G	King	Seattle	WA	98103.0	2016	AUDI	A3	Plug-in Hybrid Electric Vehicle (PHEV)	Not eligible due to low battery range	16.0	43.0	19898891	POINT	CITY OF SEATTLE - (WA) CITY OF TACOMA - (WA)	5.303300e+10	
...	
270257 1C4RJXN60R	Pierce	Joint Base Lewis McChord	WA	98433.0	2024	JEEP	WRANGLER	Plug-in Hybrid Electric Vehicle (PHEV)	Not eligible due to low battery range	21.0	28.0	266021122	POINT	PUGET SOUND ENERGY INC	5.305307e+10	
270258 1C4J3XR66N	Mason	Hoodsport	WA	98548.0	2022	JEEP	WRANGLER	Plug-in Hybrid Electric Vehicle (PHEV)	Not eligible due to low battery range	22.0	35.0	282429398	POINT	BONNEVILLE POWER ADMINISTRATION	5.304596e+10	
270259 7SAYGDEEXP	Pierce	Tacoma	WA	98406.0	2023	TESLA	MODEL Y	Battery Electric Vehicle (BEV)	Eligibility unknown as battery range has not b...	0.0	27.0	228485085	POINT	BONNEVILLE POWER ADMINISTRATION	5.305306e+10	
270260 5YJYGDDE2M	Snohomish	Bothell	WA	98021.0	2021	TESLA	MODEL Y	Battery Electric Vehicle (BEV)	Eligibility unknown as battery range has not b...	0.0	1.0	282699217	POINT	PUGET SOUND ENERGY INC	5.306105e+10	
270261 JN1BF0BA5P	Chelan	Wenatchee	WA	98801.0	2023	NISSAN	ARIYA	Battery Electric Vehicle (BEV)	Eligibility unknown as battery range has not b...	0.0	12.0	261475224	POINT	PUD NO 1 OF CHELAN COUNTY	5.300796e+10	

270262 rows × 16 columns

How can we use Linear Regression to predict the Electric Range of a vehicle?

```
In [8]: df = df[['Model Year', 'Electric Range', 'Legislative District']].dropna()
y = df['Electric Range']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
model = LinearRegression().fit(X_train, y_train)
y_pred = model.predict(X_test)

print("Coefficients: ", model.coef_)
print("Intercept: ", model.intercept_)
print("R^2: ", r2_score(y_test, y_pred))
print("Predictions: ", y_pred)

Coefficients: [-14.14570503  0.02866648]
Intercept: 28641.659603274988
R^2: 0.2958970015663
Predictions: [ 12.04246216 40.36269818 25.24323268 ... -16.22012193 -2.21774932
 39.30703825]
```

```
In [10]: df = df[['Model Year', 'Electric Range', 'Legislative District']].dropna()
y = df['Electric Range']

# Train model
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
model = LinearRegression().fit(X_train, y_train)
y_pred_train = model.predict(X_train)
y_pred_test = model.predict(X_test)

# Create visualizations
fig, axes = plt.subplots(2, 2, figsize=(15, 12))
fig.suptitle('Linear Regression: Predicting Electric Range', fontsize=16)

# 1. Actual vs Predicted (Training)
axes[0,0].scatter(y_train, y_pred_train, alpha=0.7, color='blue')
axes[0,0].plot([y_train.min(), y_train.max()], [y_train.min(), y_train.max()], 'r--', lw=2)
axes[0,0].set_xlabel('Actual Electric Range')
axes[0,0].set_ylabel('Predicted Electric Range')
axes[0,0].set_title('Training: R^2 = %r' % (r2_score(y_train, y_pred_train),))

# 2. Actual vs Predicted (Testing)
axes[0,1].scatter(y_test, y_pred_test, alpha=0.7, color='green')
axes[0,1].plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], 'r--', lw=2)
axes[0,1].set_xlabel('Actual Electric Range')
axes[0,1].set_ylabel('Predicted Electric Range')
axes[0,1].set_title('Testing: R^2 = %r' % (r2_score(y_test, y_pred_test),))

# 3. Residuals Plot
residuals = y_test - y_pred_test
axes[1,0].scatter(y_test, residuals, alpha=0.7, color='orange')
axes[1,0].plot([-100, 100], [0, 0], 'r--', linestyle='--')
axes[1,0].set_xlabel('Actual Electric Range')
axes[1,0].set_ylabel('Residuals')
axes[1,0].set_title('Testing: R^2 = %r' % (r2_score(y_test, y_pred_test),))

# 4. Model Year vs Electric Range with regression line
scatters = axes[1,1].scatter(df['Model Year'], df['Electric Range'],
                           c=df['Legislative District'], s=100, alpha=0.7)
model_year_line = axes[1,1].scatter(df['Model Year'], df['Electric Range'],
                                    c=df['Legislative District'], s=100, alpha=0.7)
model = LinearRegression().fit(df['Model Year'], df['Electric Range'])
model_year_line = axes[1,1].plot(df['Model Year'], df['Electric Range'],
                                 c=df['Legislative District'], s=100, alpha=0.7)
axes[1,1].set_xlabel('Model Year')
axes[1,1].set_ylabel('Electric Range')
axes[1,1].set_title('Model Year vs Electric Range')
plt.colorbar(scatters, ax=axes[1,1], label='Legislative District')
axes[1,1].legend()

plt.tight_layout()
plt.show()
```



C:\Users\LENOVO\anaconda3\lib\site-packages\c:\Python\3.7\lib\site-packages\ipython\core\pylabtools.py:170: UserWarning: Creating legend with loc="best" can be slow with large amounts of data.
fig.canvas.print_figure(options_to, **kw)

Linear Regression: Predicting Electric Range

Training: R² = 0.297

Testing: R² = 0.296

Actual Electric Range

Predicted Electric Range

Residuals Plot

Model Year vs Electric Range

Model Coefficients: Model Year=-14.15, District=0.03

Intercept: 28641.66

R²: 0.296

RMSE: 66.52

What independent variables (features) can be used to predict Electric Range? (e.g., Model Year, Base MSRP, Make)

```
In [11]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
from sklearn.preprocessing import LabelEncoder

# Potential numeric features for 'Electric Range prediction'
numeric_features = ['Model Year', 'Legislative District', 'Postal Code']

# Create feature matrix (encode categorical, handle missing values)
df_numeric = df[numeric_features + ['Electric Range']].dropna()
X_numeric = df_numeric[numeric_features]

# Encode categorical features (one-hot for small dataset)
df_cat = pd.get_dummies(df[categorical_features + ['Electric Range']].dropna(), drop_first=True)
X_cat = df_cat.drop(['Electric Range'], axis=1)

print("Available Numeric Features: ", numeric_features)
print("Available Categorical Features: ", categorical_features)
print("Dataset shape after cleaning: ", df.shape)
```

How do we handle categorical variables like Make and Model in regression analysis?

```
In [14]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import r2_score

df_clean = df[['Model Year', 'Make', 'Model', 'Electric Range', 'Model Year']].dropna()

# One-hot encoding (make)
X_cat = df_clean[['Model Year', 'Model']]
y = df_clean[['Electric Range']]

che_b = OneHotEncoder(drop='first', sparse_output=False) # drop='first' avoids multicollinearity
X_cat = che_b.fit_transform(X_cat)
X_cat = che_b.named_steps['onehotencoder'].get_feature_names_out()
X_cat_df = pd.DataFrame(X_cat, columns=che_b.get_feature_names_out())

OneHotEncoder: Model Year=0.0000000000000002, Model=0.0000000000000002, Model Year=0.0000000000000002
```

Sample columns: ['Make_ALFA ROMEO', 'Make_AUDI', 'Make_AURUS DYNAMICS', 'Make_BENTLEY', 'Make_BMW']

Model Coefficients: Model Year=-14.15, Make=0.0000000000000002, Model=0.0000000000000002

Intercept: 28641.659603274988

R²: 0.2958970015663

Predictions: [12.04246216 40.36269818 25.24323268 ... -16.22012193 -2.21774932
 39.30703825]

What is the R² score of the model, and what does it indicate about prediction accuracy?

```
In [18]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score

df1 = df[['Model Year', 'Electric Range', 'Legislative District']].dropna()
y = df1['Electric Range']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
model = LinearRegression().fit(X_train, y_train)
y_pred = model.predict(X_test)

print("R^2 Score: (%r)" % (r2_score(y_test, y_pred),))
print("Test R^2: (%r)" % (r2_score(y_test, y_pred),))
print("RMSE: (%r)" % (np.sqrt(mean_squared_error(y_test, y_pred)),))
print("MSE: (%r)" % (mean_squared_error(y_test, y_pred),))

C:\Users\LENOVO\anaconda3\lib\site-packages\c:\Python\3.7\lib\site-packages\ipython\core\pylabtools.py:170: UserWarning: Creating legend with loc="best" can be slow with large amounts of data.
fig.canvas.print_figure(options_to, **kw)
```

What steps are needed to improve the accuracy of the Linear Regression model?

```
In [20]: # Add categorical features via one-hot encoding
df_cat = df[['Model Year', 'Electric Vehicle Type', 'Electric Range']].dropna()

# One-hot encode Make and Vehicle Type
X_cat = pd.get_dummies(df_cat[['Model Year', 'Make', 'Electric Vehicle Type']], drop_first=True)
y = df_cat[['Electric Range']]

print("Available Categorical Features: ['Model Year', 'Electric Vehicle Type']")
print("Dataset shape after cleaning: (%r, %r)" % (df_cat.shape, df_cat.shape))

# Enhanced feature matrix
X_enhanced = pd.concat([df[['Model Year', 'Postal Code', 'Clear Alternative Fuel Vehicle (CAVF) Eligibility']], df_cat], axis=1)

# Create feature matrix (encode categorical, handle missing values)
df_numeric = df[numeric_features + ['Electric Range']].dropna()
X_numeric = df_numeric[numeric_features]

# Encode categorical features (one-hot for small dataset)
df_cat = pd.get_dummies(df[categorical_features + ['Electric Range']].dropna(), drop_first=True)
X_cat = df_cat.drop(['Electric Range'], axis=1)

print("Available Numeric Features: ", numeric_features)
print("Available Categorical Features: ", categorical_features)
print("Dataset shape after cleaning: (%r, %r)" % (df.shape, df.shape))
```

How do we handle categorical variables like Make and Model in regression analysis?

```
In [17]: # Combine with numeric features
X_numeric = df_clean[['Model Year']]
X_final = pd.concat([X_numeric.reset_index(drop=True), X_cat.reset_index(drop=True)], axis=1)

# Train-test split and model
X_train, X_test, y_train, y_test = train_test_split(X_final, y.reset_index(drop=True), test_size=0.3, random_state=42)
model = LinearRegression().fit(X_train, y_train)
y_pred = model.predict(X_test)

print("R^2 Score: (%r)" % (r2_score(y_test, y_pred),))
print("Test R^2: (%r)" % (r2_score(y_test, y_pred),))
print("RMSE: (%r)" % (np.sqrt(mean_squared_error(y_test, y_pred)),))
print("MSE: (%r)" % (mean_squared_error(y_test, y_pred),))

C:\Users\LENOVO\anaconda3\lib\site-packages\c:\Python\3.7\lib\site-packages\ipython\core\pylabtools.py:170: UserWarning: Creating legend with loc="best" can be slow with large amounts of data.
fig.canvas.print_figure(options_to, **kw)
```

Model Coefficients: Model Year=-14.15, District=0.03

Intercept: 28641.66

R²: 0.296

RMSE: 66.52

What independent variables (features) can be used to predict Electric Range? (e.g., Model Year, Base MSRP, Make)

```
In [11]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
from sklearn.preprocessing import LabelEncoder

# Potential numeric features for 'Electric Range prediction'
numeric_features = ['Model Year', 'Legislative District', 'Postal Code']

# Create feature matrix (encode categorical, handle missing values)
df_numeric = df[numeric_features + ['Electric Range']].dropna()
X_numeric = df_numeric[numeric_features]

# Encode categorical features (one-hot for small dataset)
df_cat = pd.get_dummies(df[categorical_features + ['Electric Range']].dropna(), drop_first=True)
X_cat = df_cat.drop(['Electric Range'], axis=1)

print("Available Numeric Features: ", numeric_features)
print("Available Categorical Features: ", categorical_features)
print("Dataset shape after cleaning: (%r, %r)" % (df.shape, df.shape))
```

How do we handle categorical variables like Make and Model in regression analysis?

```
In [14]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import r2_score

df_clean = df[['Model Year', 'Make', 'Model', 'Electric Range', 'Model Year']].dropna()

# One-hot encoding (make)
X_cat = df_clean[['Model Year', 'Model']]
y = df_clean[['Electric Range']]

che_b = OneHotEncoder(drop='first', sparse_output=False) # drop='first' avoids multicollinearity
X_cat = che_b.fit_transform(X_cat)
X_cat = che_b.named_steps['onehotencoder'].get_feature_names_out()
X_cat_df = pd.DataFrame(X_cat, columns=che_b.get_feature_names_out())

OneHotEncoder: Model Year=0.0000000000000002, Model=0.0000000000000002, Model Year=0.0000000000000002
```

Sample columns: ['Make_ALFA ROMEO', 'Make_AUDI', 'Make_AURUS DYNAMICS', 'Make_BENTLEY', 'Make_BMW']

Model Coefficients: Model Year=-14.15, Make=0.0000000000000002, Model=0.0000000000000002

Intercept: 28641.659603274988

R²: 0.2958970015663

Predictions: [12.04246216 40.36269818 25.24323268 ... -16.22012193 -2.21774932
 39.30703825]

What is the R² score of the model, and what does it indicate about prediction accuracy?

```
In [18]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score

df1 = df[['Model Year', 'Electric Range', 'Legislative District']].dropna()
y = df1['Electric Range']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
model = LinearRegression().fit(X_train, y_train)
y_pred = model.predict(X_test)

print("R^2 Score: (%r)" % (r2_score(y_test, y_pred),))
print("Test R^2: (%r)" % (r2_score(y_test, y_pred),))
print("RMSE: (%r)" % (np.sqrt(mean_squared_error(y_test, y_pred)),))
print("MSE: (%r)" % (mean_squared_error(y_test, y_pred),))

C:\Users\LENOVO\anaconda3\lib\site-packages\c:\Python\3.7\lib\site-packages\ipython\core\pylabtools.py:170: UserWarning: Creating legend with loc="best" can be slow with large amounts of data.
fig.canvas.print_figure(options_to, **kw)
```

What steps are needed to improve the accuracy of the Linear Regression model?

```
In [20]: # Add categorical features via one-hot encoding
df_cat = df[['Model Year', 'Electric Vehicle Type', 'Electric Range']].dropna()

# One-hot encode Make and Vehicle Type
X_cat = pd.get_dummies(df_cat[['Model Year', 'Make', 'Electric Vehicle Type']], drop_first=True)
y = df_cat[['Electric Range']]

print("Available Categorical Features: ['Model Year', 'Electric Vehicle Type']")
print("Dataset shape after cleaning: (%r, %r)" % (df_cat.shape, df_cat.shape))

# Enhanced feature matrix
X_enhanced = pd.concat([df[['Model Year', 'Postal Code', 'Clear Alternative Fuel Vehicle (CAVF) Elig
```